



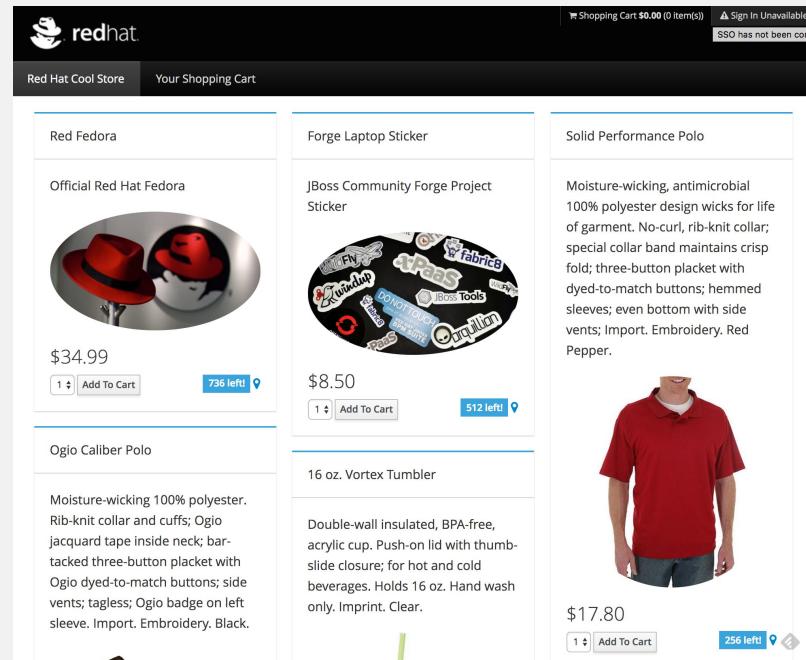
# MULTI PRODUCT MICROSERVICES LAB

James Falkner  
Sr. Technical Marketing Manager, Middleware  
April 2017

# COOLSTORE APPLICATION DEMO

# COOLSTORE APPLICATION

- Online shop for selling products
- Web-based polyglot application using
  - AngularJS
  - Node
  - Java with multiple frameworks (Spring, Java EE, etc)
- Microservices architecture
- Deployed using containers



# COOLSTORE SERVICES

Red Hat Cool Store   Your Shopping Cart

Red Fedora

Official Red Hat Fedora



\$34.99

1 ↑ Add To Cart

736 left! ⏭

Forge Laptop Sticker

JBoss Community Forge Project Sticker



\$8.50

1 ↑ Add To Cart

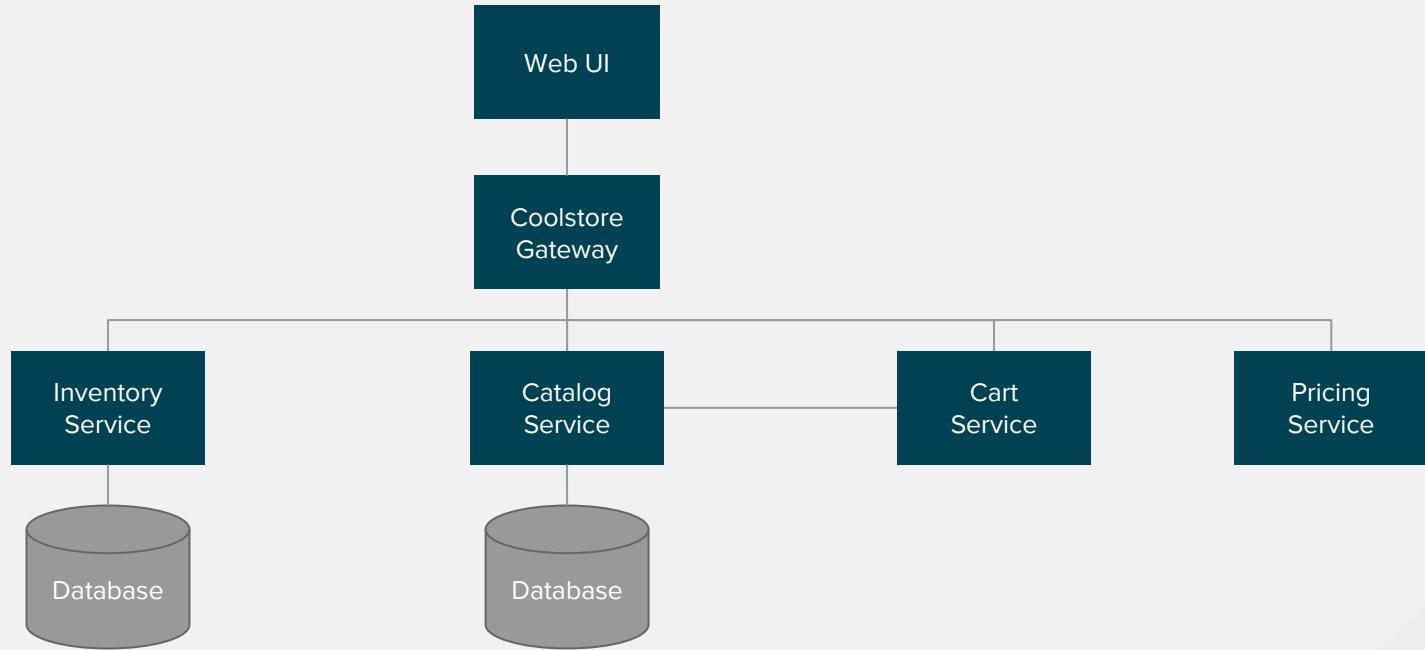
512 left! ⏭

Solid Performance Polo

Moisture-wicking, antimicrobial 100% polyester design wicks for life of garment. No-curl, rib-knit collar; special collar band maintains crisp fold; three-button placket with dyed-to-match buttons; hemmed sleeves; even bottom with side vents; Import. Embroidery. Red Pepper.



# COOLSTORE SERVICES



# DECK GUIDE: ORDERING THE DEMO

1. Log in into RHPDS (demo system): <https://rhpds.redhat.com>
2. Order the demo

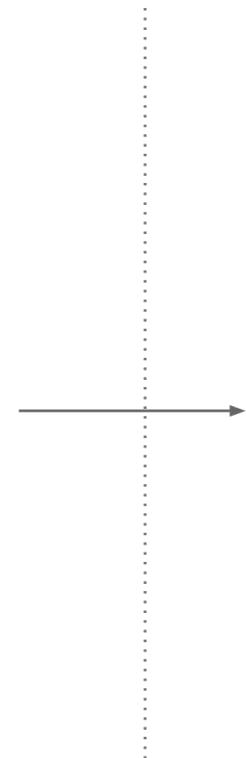
The screenshot shows the Red Hat CloudForms Management Engine interface. On the left, there's a sidebar with navigation icons and a 'Service Catalogs' section. Under 'All Services', several categories are listed: Cloud Infrastructure Demos, Community Cloud Demo, Datacenter Infrastructure Demos, Middleware Solutions Demos, and Multi-product Demos. The 'Modem Application Architecture Demo' item is selected, highlighted with a blue background. To the right, detailed information about this demo is displayed. The service name is "Modern Application Architecture Demo". The 'Name' field shows "Modern Application Architecture Demo" and the 'Description' field shows "Modern Application Architecture Demo". The 'Long Description' section contains a detailed description of the demo, mentioning it's a polyglot microservices application using JBoss Middleware, NodeJS, Spring Boot, and more, deployed on OpenShift. It highlights scenarios like containerized microservices, agile integration, service resilience, fault tolerance, developer productivity, CI/CD, etc. Below this, provisioning details are provided: 'Provisioning Time: ~10 min', 'Demo Guide: <http://guide-ci-USERNAME.cloudapps.na.openshift.opentlc.com>', and a note to replace 'USERNAME' with the user's RH username. At the bottom right of the main content area is a 'Order' button.

3. Demo will be ready ~15 min later

# MODULES VS. USE CASES

## DEMO MODULES

DEMO MODULE A
DEMO MODULE B
DEMO MODULE C
DEMO MODULE D
DEMO MODULE E
DEMO MODULE F
DEMO MODULE G
DEMO MODULE H
DEMO MODULE I
DEMO MODULE J



## DEMO USE-CASES

### Demo: Use-Case 1

DEMO MODULE A
DEMO MODULE B
DEMO MODULE D
DEMO MODULE E

### Demo: Use-Case 3

DEMO MODULE D
DEMO MODULE H
DEMO MODULE I
DEMO MODULE J

### Demo: Use-Case 2

DEMO MODULE I
DEMO MODULE E
DEMO MODULE F

### Demo: Use-Case 4

DEMO MODULE B
DEMO MODULE F
DEMO MODULE G

# GUIDES AND VIDEOS

These demos are targeted at Red Hat Partners and SAs, SSAs and SSPs and provide a set of deliverables to facilitate the demo experience of reach target group. The following is provided for each demo module in order to train the presenter on how to deliver the demo:

- **Guide**: describes how to demo, what to do and what to explain during the demo
- **Slide**: supporting slides to be used during the demo for explaining concepts
- **Video**: recording of the live demo using the guides and slides
- **Simulation**: an interactive recorded demo which can be used to demo the scenario or train for a live-demo
- **Scripts**: in order to setup the demo infrastructure on any OpenShift environment with enough resources and quota
- **Environment**: automatic provisioning of the demo infrastructure on [Red Hat Product Demo System](#) at <https://rhpds.redhat.com>

# GUIDES AND VIDEOS

## All Modules

1. CoolStore Demo Application Overview

2. Containerized Microservices

3. Agile Integration for Microservices

4. Service Resilience and Fault Tolerance

5. Creating On-Demand Environment

6. Increase Developer Productivity with Containers (with JBoss Developer Studio)

7. Building Quality into Development Process (with JBoss Developer Studio)

8. Building Quality into Development Process (without JBoss Developer Studio)

9. Increasing Delivery Speed with CI/CD

10. Continuous Security Compliance for Containers

## Workshopper

[Source code @ Github](#)

All Modules

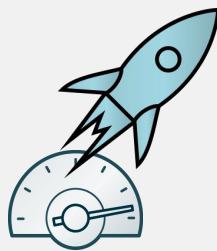
Multi-product Demo: Quick Code-to-Production without Downtime

Multi-product Demo: Agile Integration

Multi-product Demo: Modern App Architecture

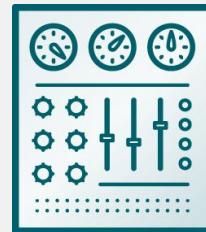
# MICROSERVICES ARCHITECTURE OVERVIEW

# VALUES OF MICROSERVICES



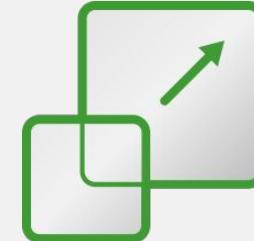
FAST TIME TO MARKET

Small autonomous services can be developed and delivered faster



EFFICIENCY

Automating delivery and monitoring of small services is easier



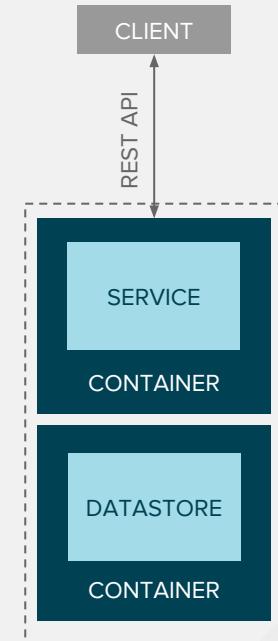
SCALABILITY

Fine grained scalability is easier and uses less resources

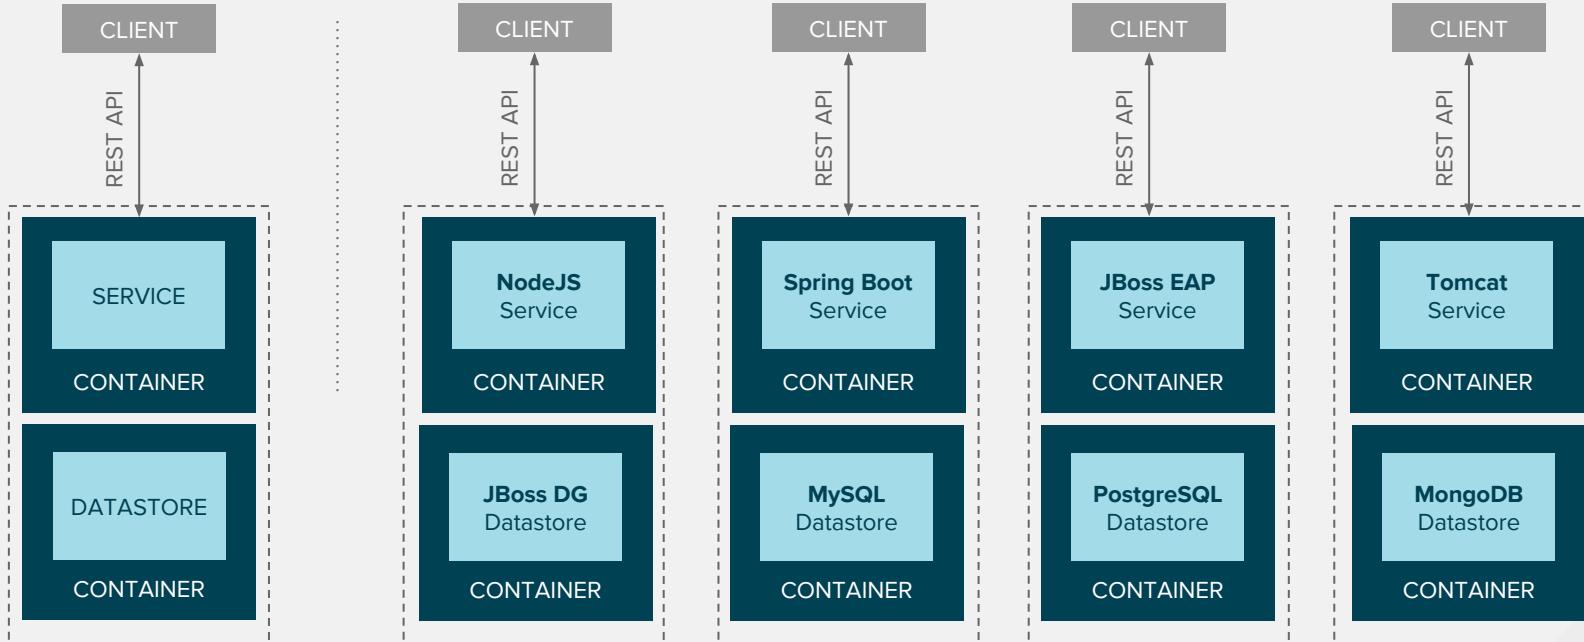
# MICROSERVICES ARCHITECTURE

Principles of microservices architecture:

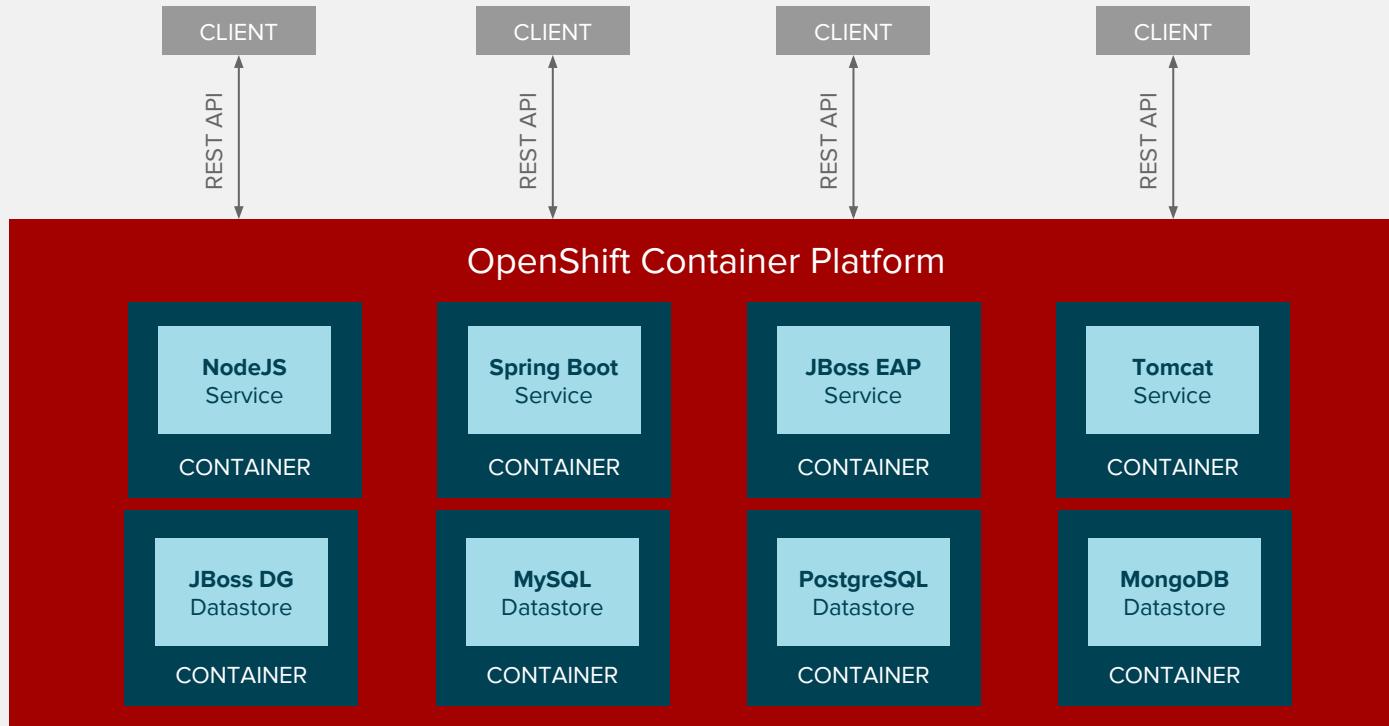
- Model around Business Domain
- Culture of Automation (DevOps)
- Deploy Independently
- Scale Independently
- Release Independently
- Combined to form a system or application
- Antifragile - increased robustness and resilience under pressure
- Polyglot (language and framework independence)
- API / contract focused
- Decentralized data management



# MICROSERVICES ARCHITECTURE



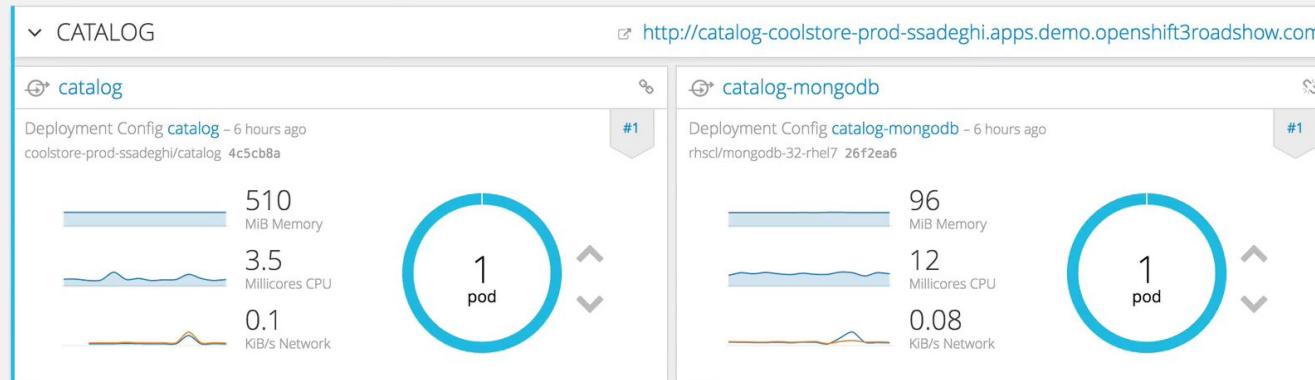
# MICROSERVICES ARCHITECTURE



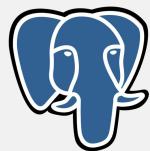
# CONTAINER DEPLOYMENT & MANAGEMENT

OpenShift Container Platform provides features like:

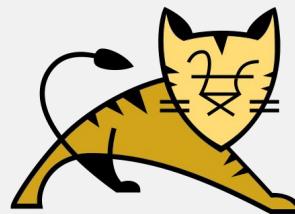
- Deploy Independently
- Scale Independently
- Release Independently
- CI/CD (DevOps)
- Antifragile
- Polyglot



# POLYGLOT



PostgreSQL



RED HAT® JBOSS®  
FUSE

RED HAT® JBOSS®  
ENTERPRISE  
APPLICATION PLATFORM



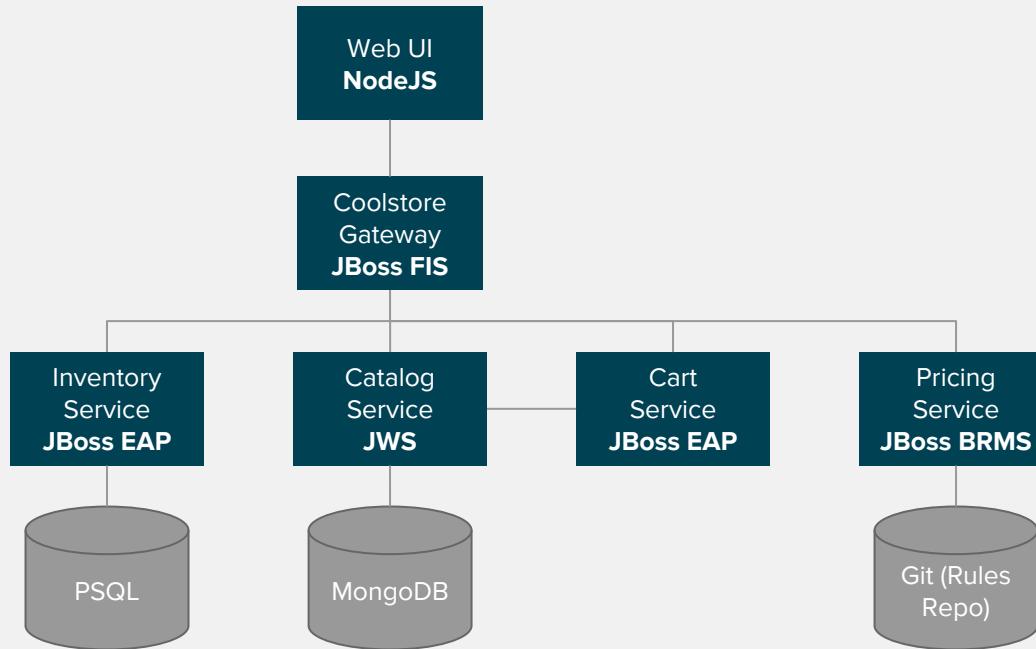
WildFly



RED HAT® JBOSS®  
BRMS



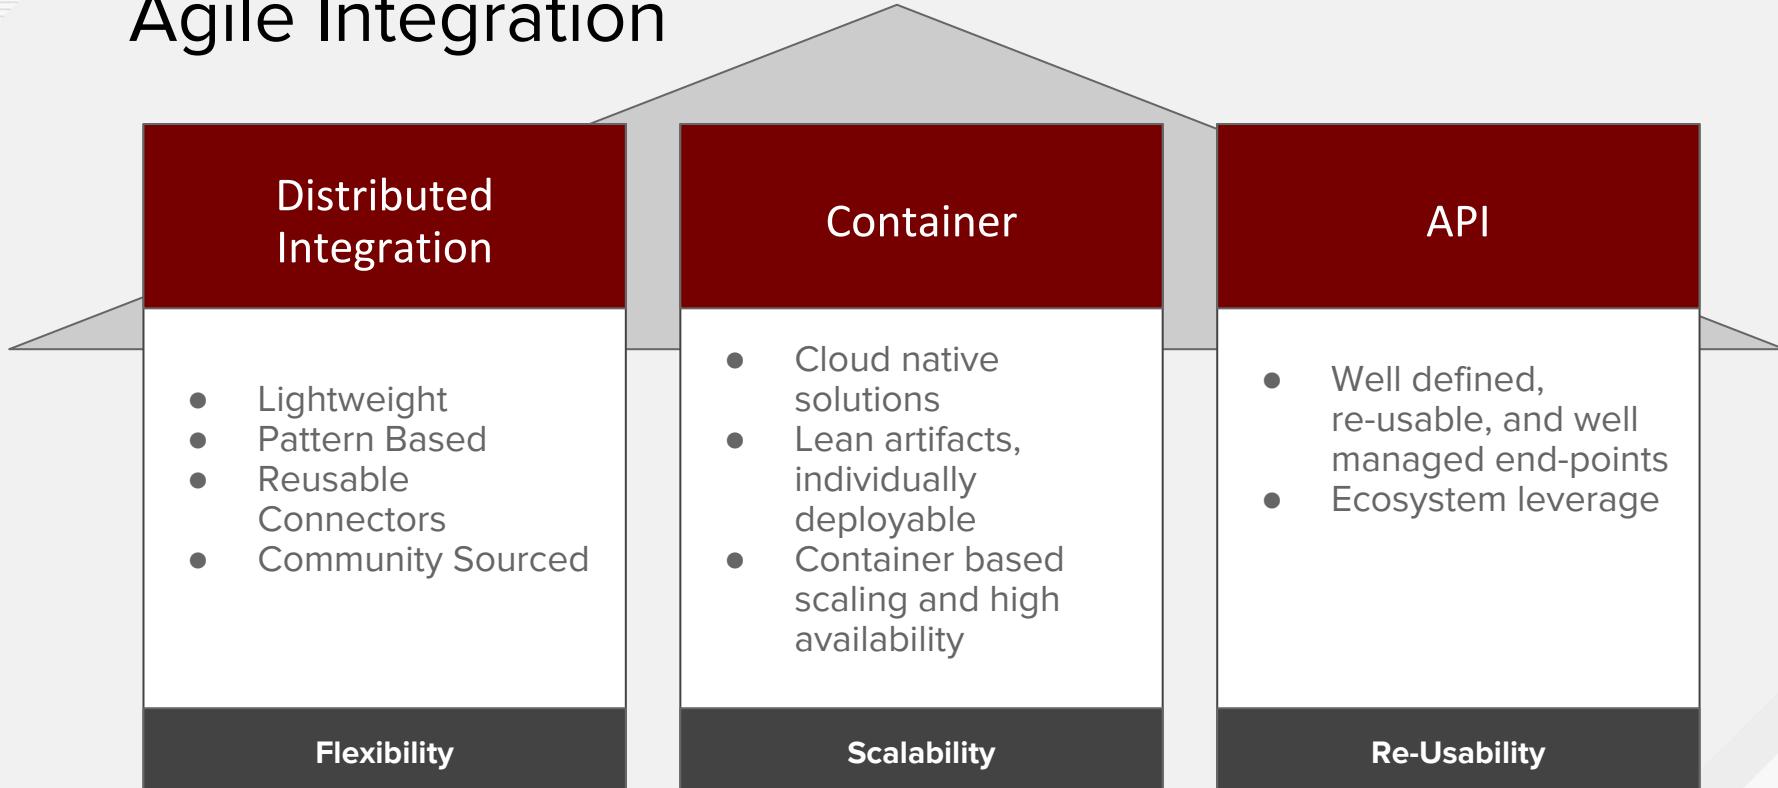
# DEMO APPLICATION



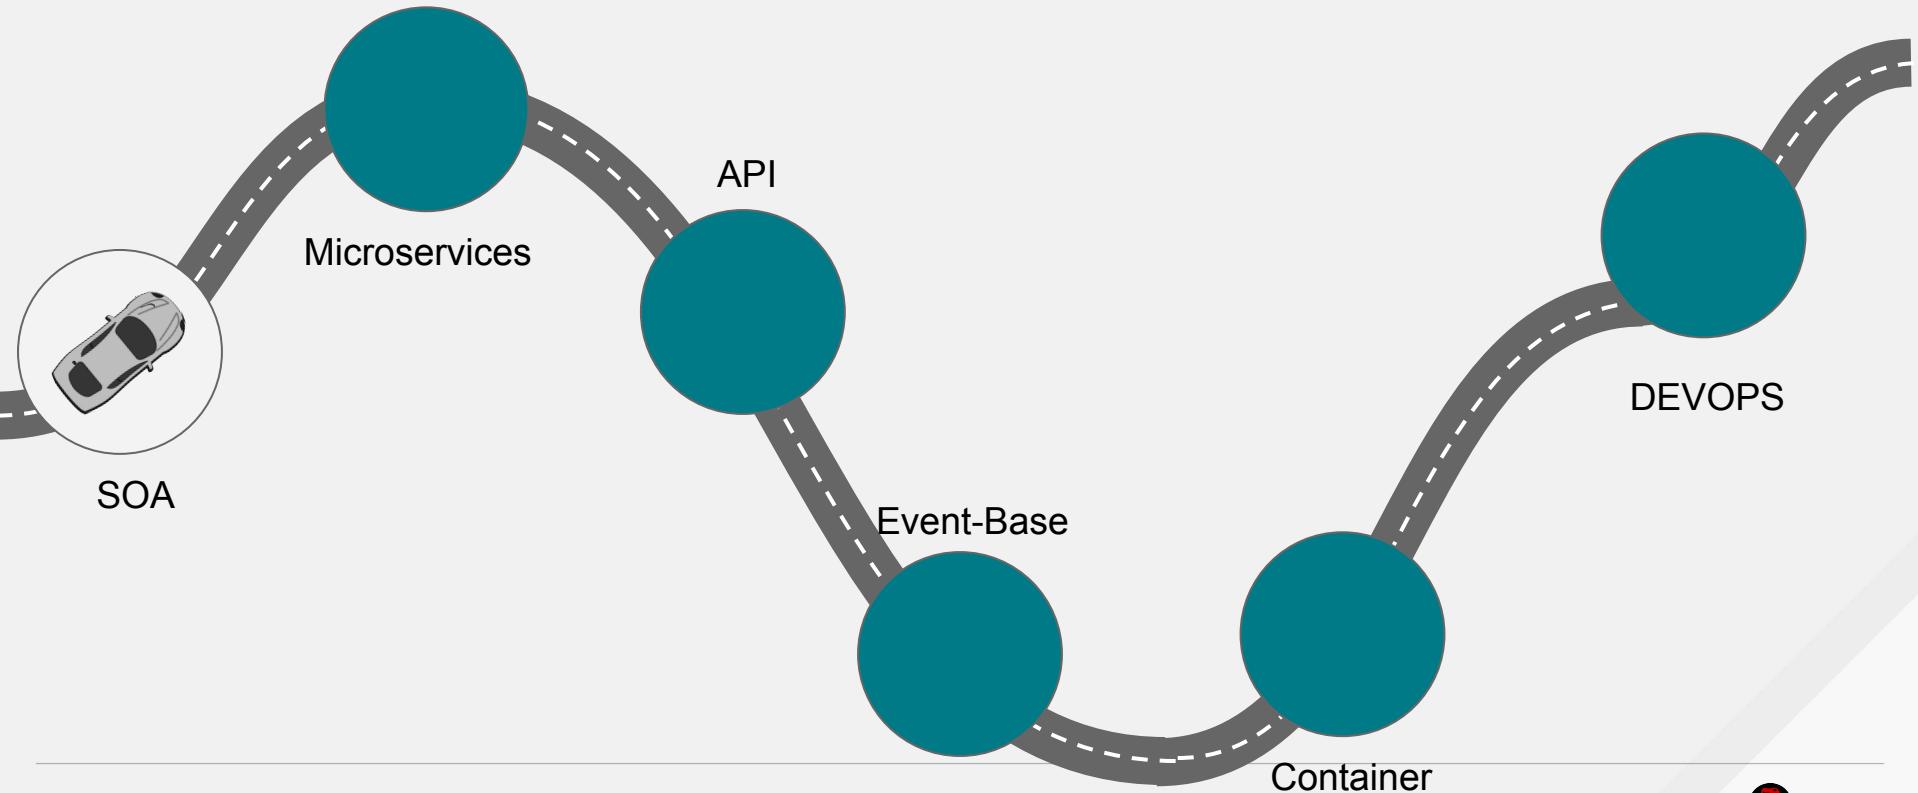
The background of the slide features a dark teal overlay with a subtle, abstract architectural pattern. It depicts several modern skyscrapers with glass facades and steel frames, some with vertical cladding. The perspective is from a low angle, looking up at the buildings against a lighter sky.

# AGILE INTEGRATION FOR MICROSERVICES

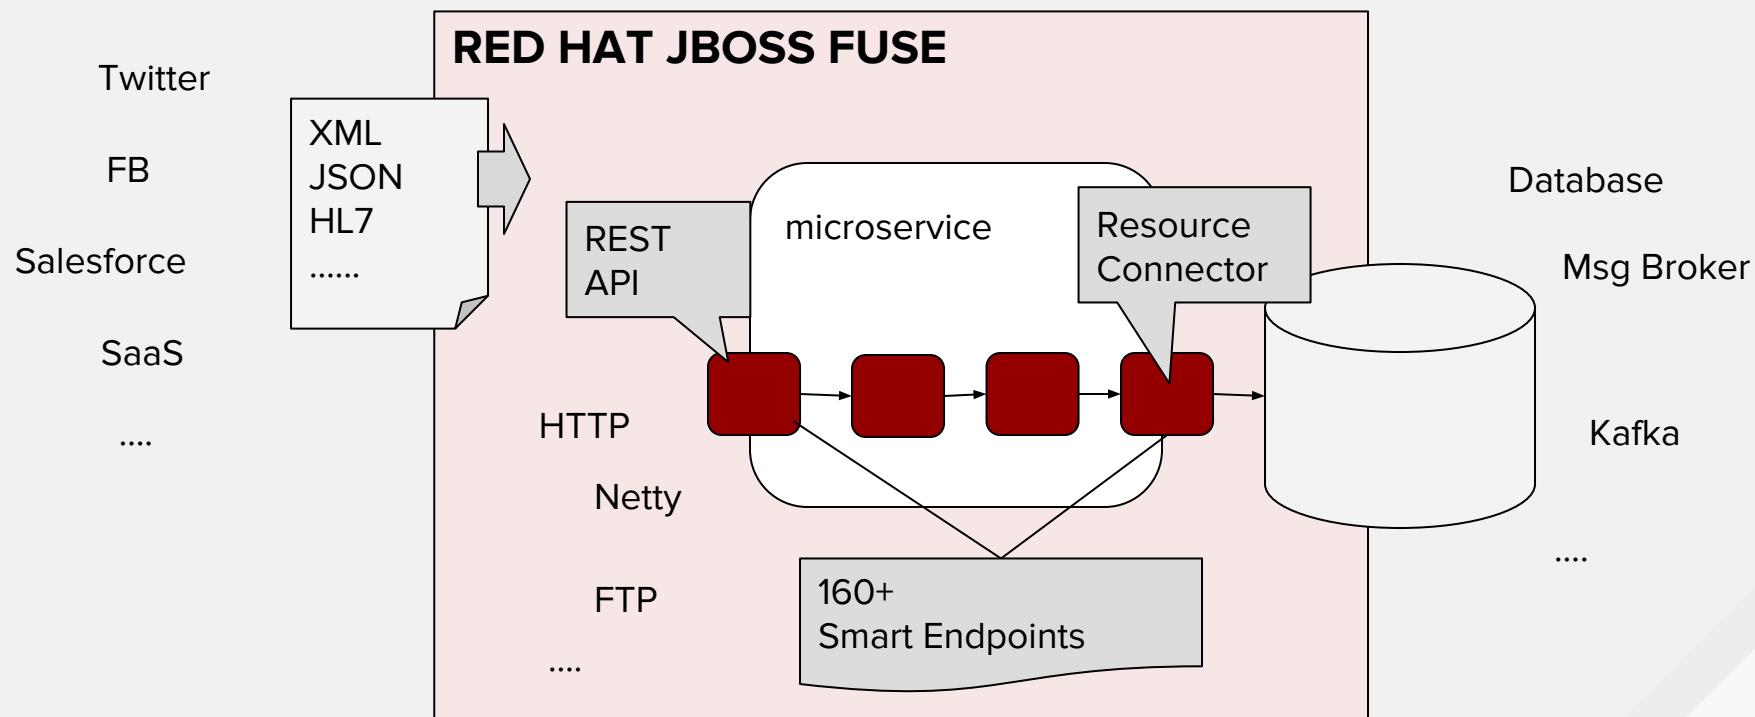
# Agile Integration



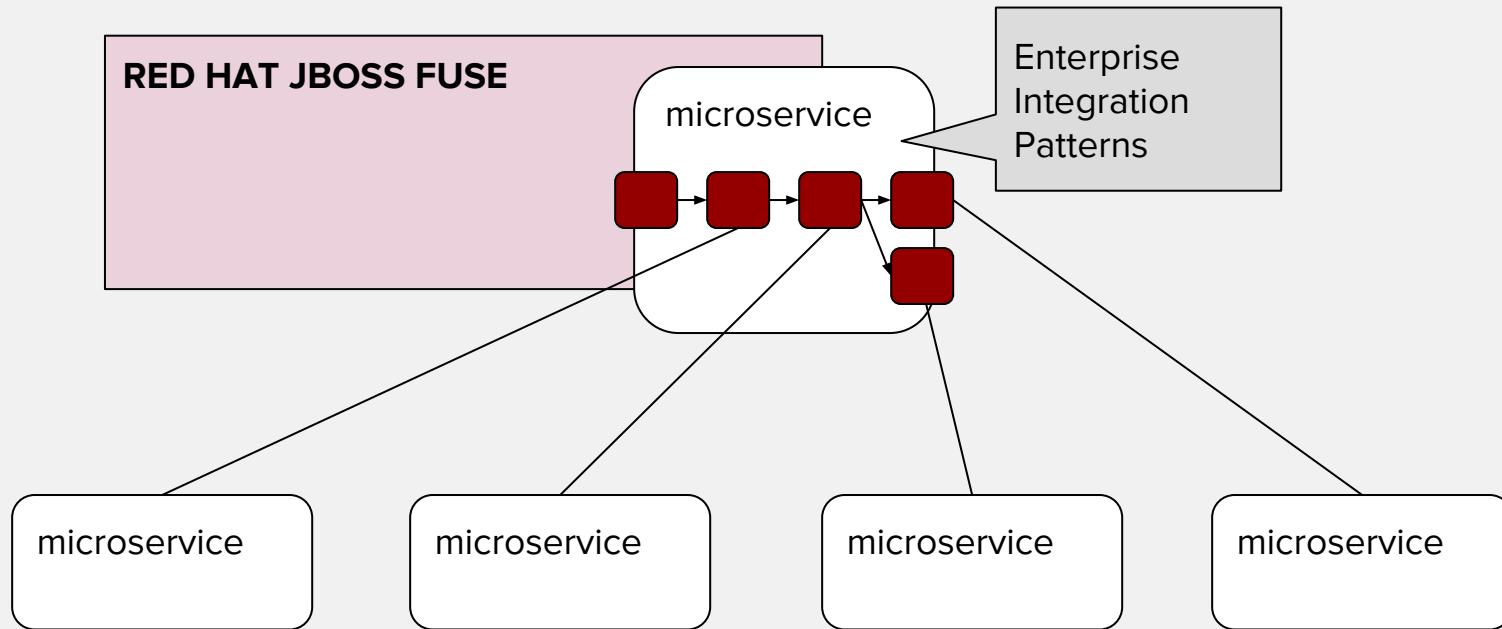
# AGILE INTEGRATION



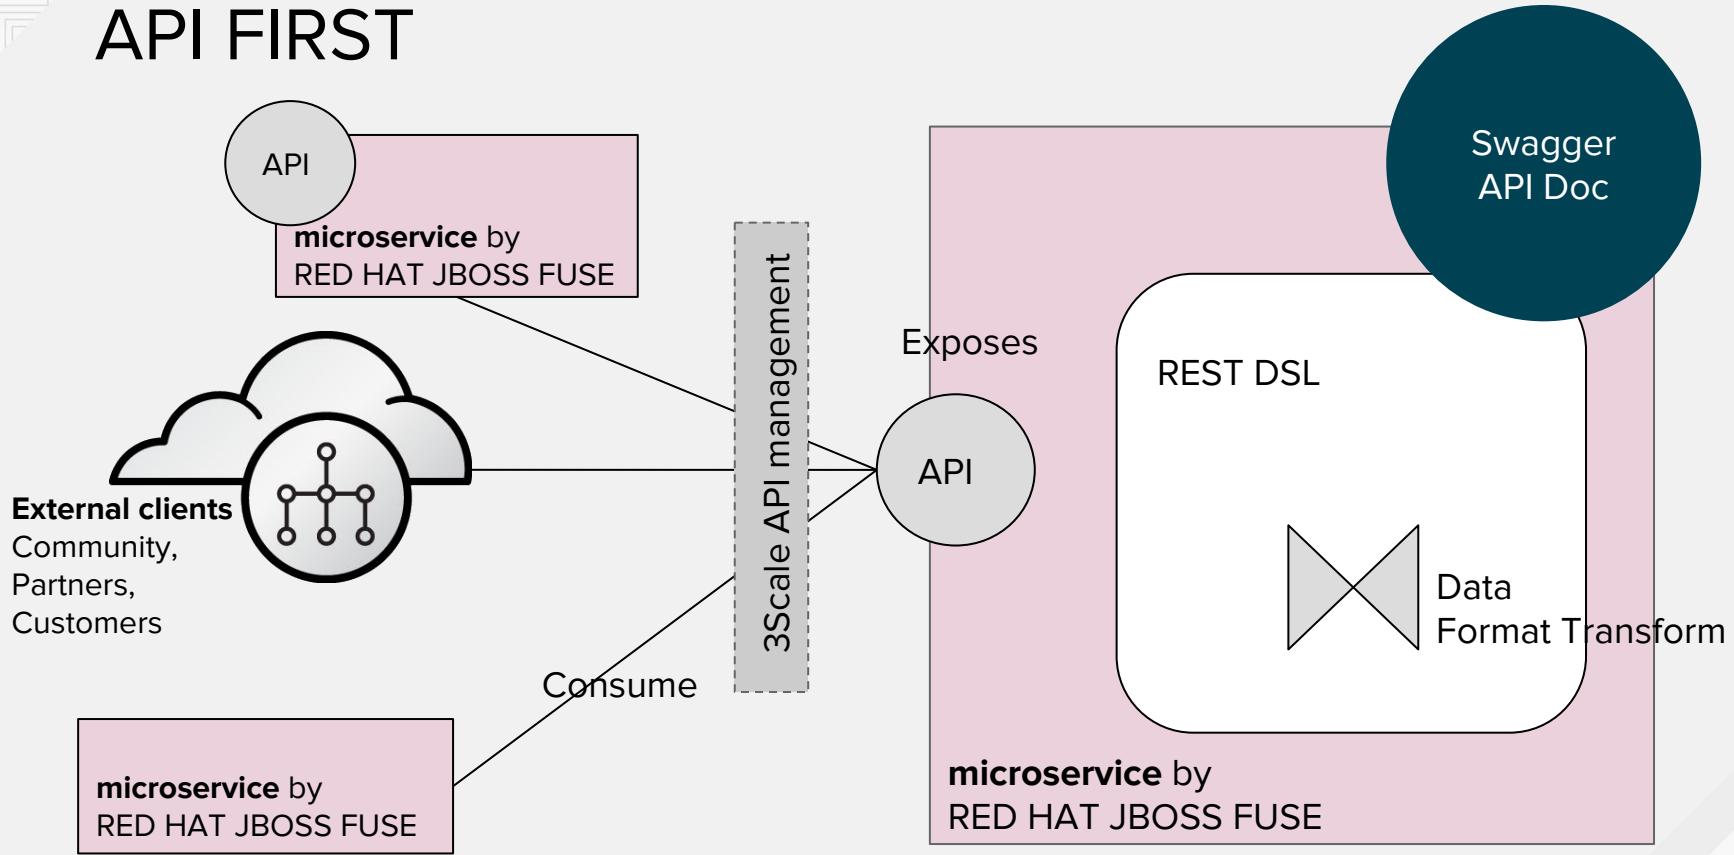
# BUILDING MICROSERVICES



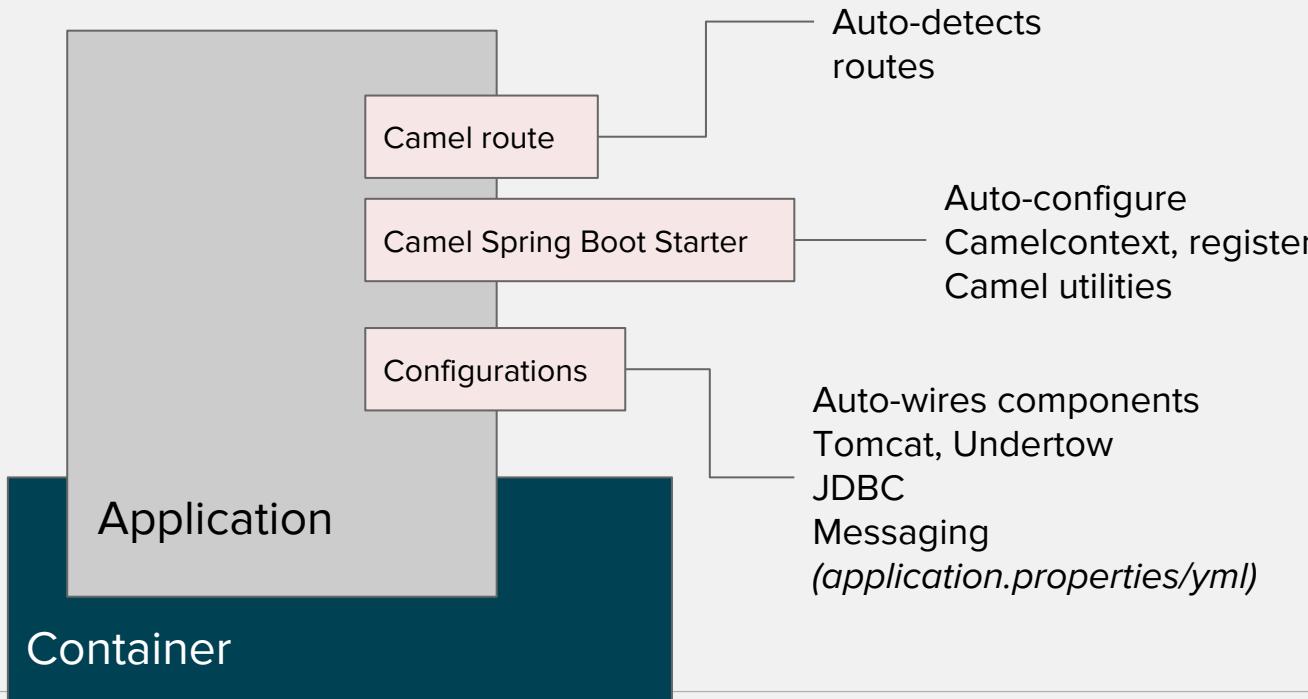
# COMPOSING MICROSERVICES



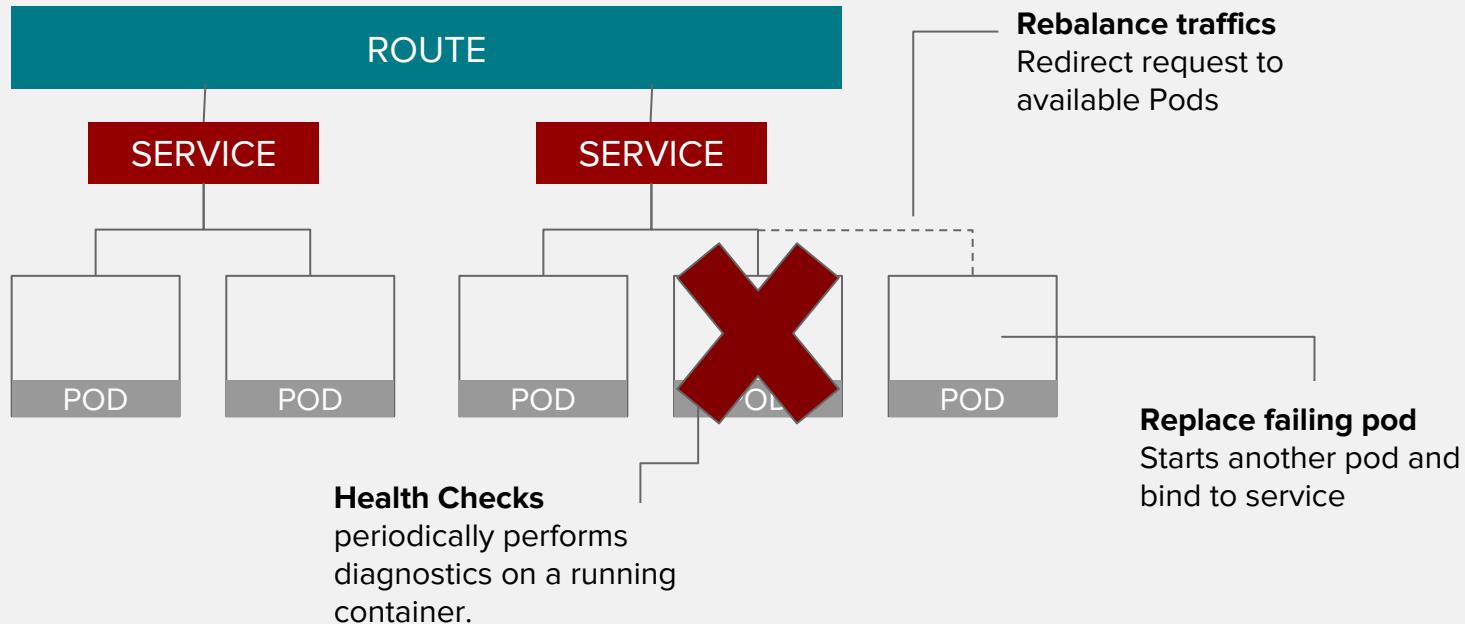
# API FIRST



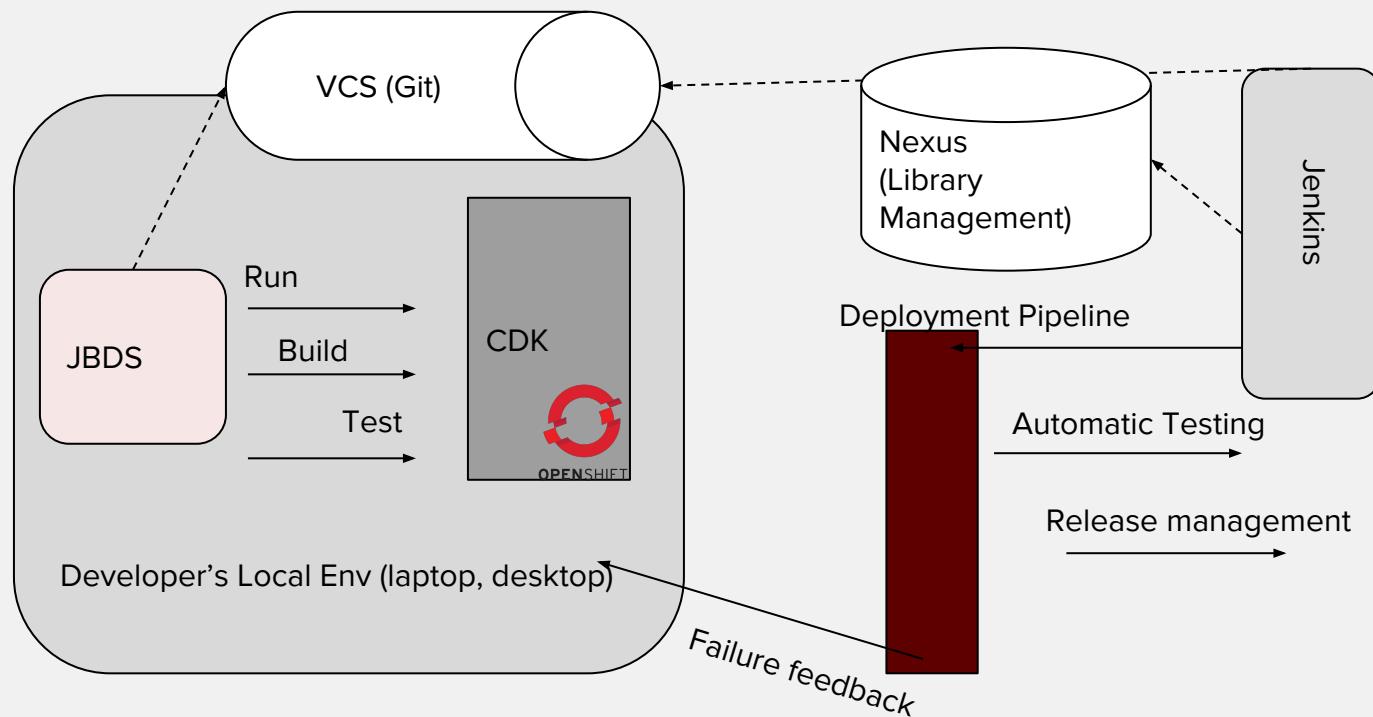
# Fuse and Spring Boot



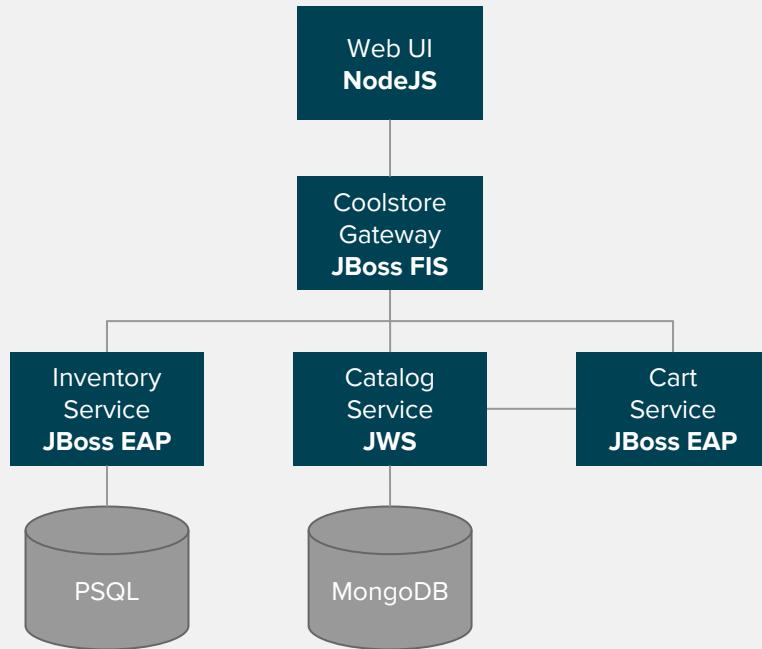
# FAILURE RECOVERY



# CONTAINER



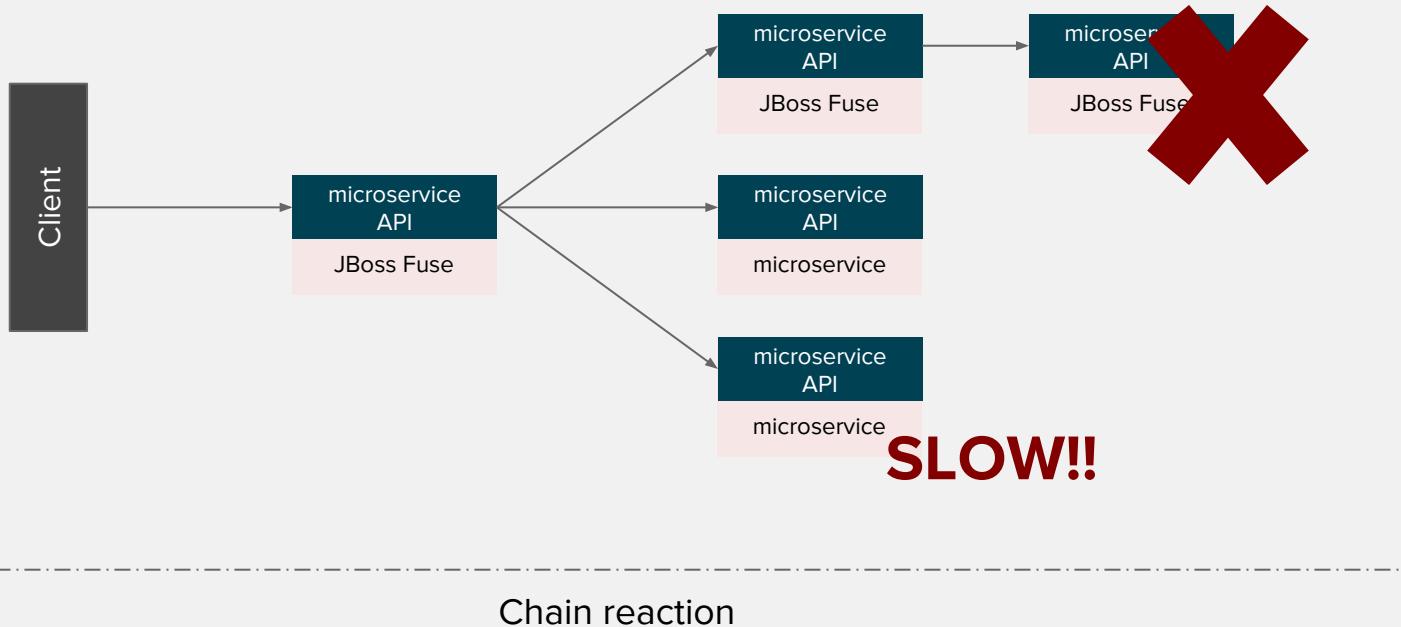
# DEMO APPLICATION



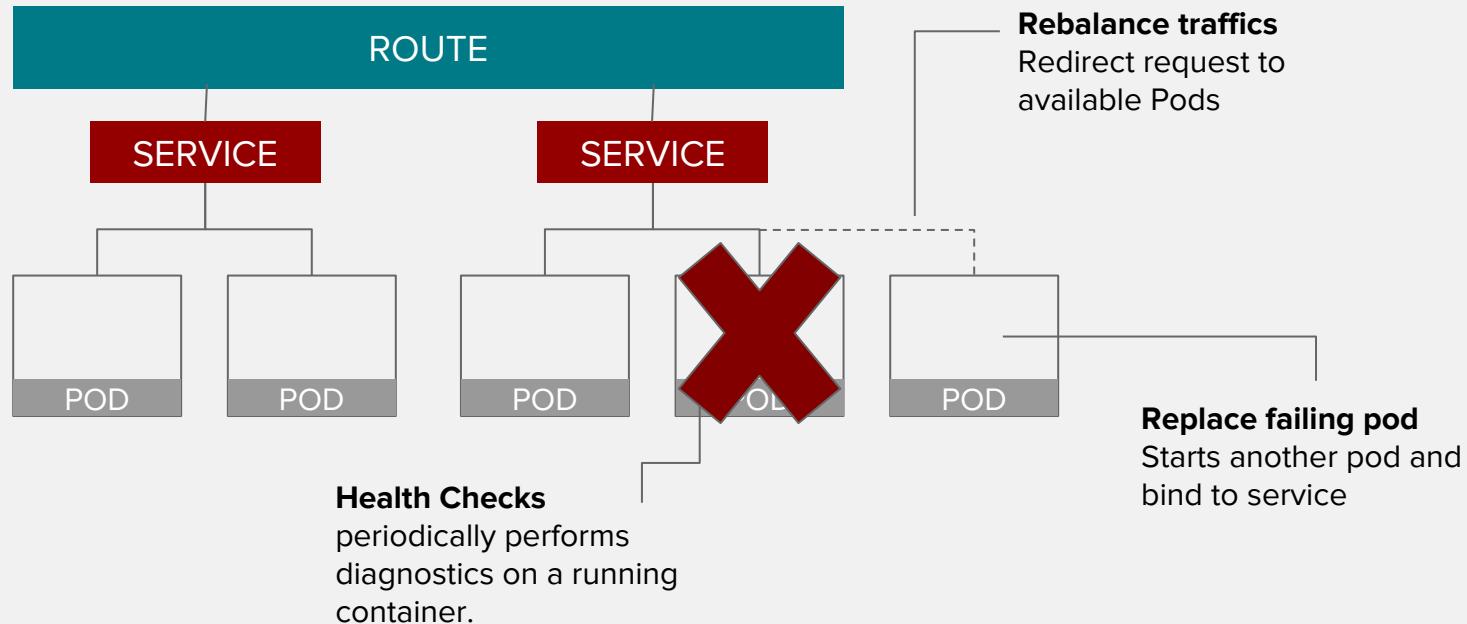


# SERVICE RESILIENCE AND FAULT TOLERANCE

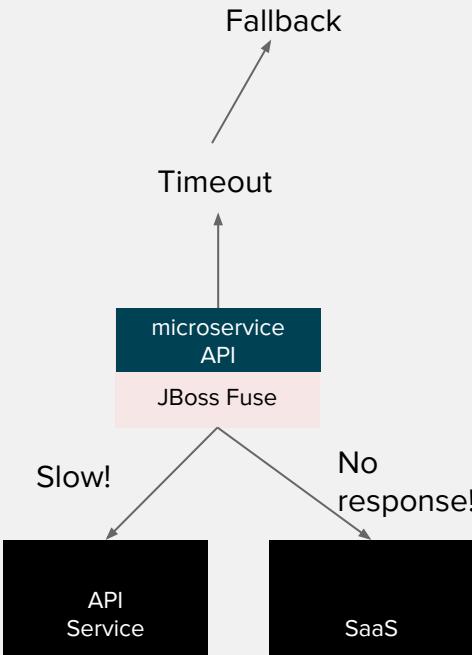
# SERVICE RESILIENCE



# FAULT TOLERANCE

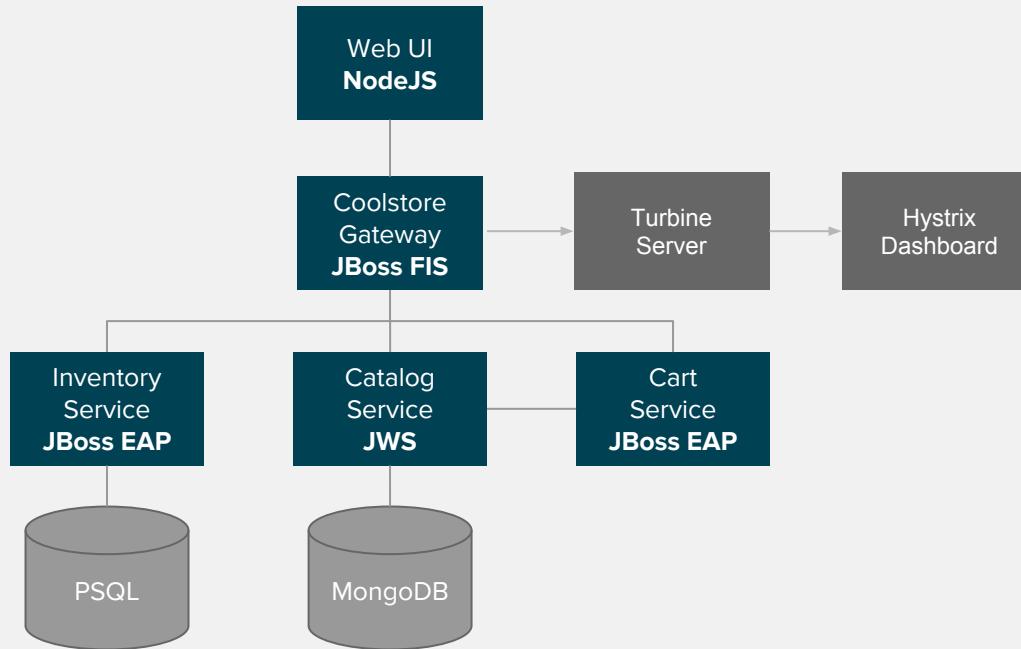


# Hybris EIP



```
<camelContext xmlns="http://camel.apache.org/schema/spring">
  <route>
    <from uri="direct:start"/>
    <hybris>
      <to uri="http://fooservice.com/slow"/>
      <onFallback>
        <transform>
          <constant>Fallback message</constant>
        </transform>
      </onFallback>
    </hybris>
    <to uri="mock:result"/>
  </route>
</camelContext>
```

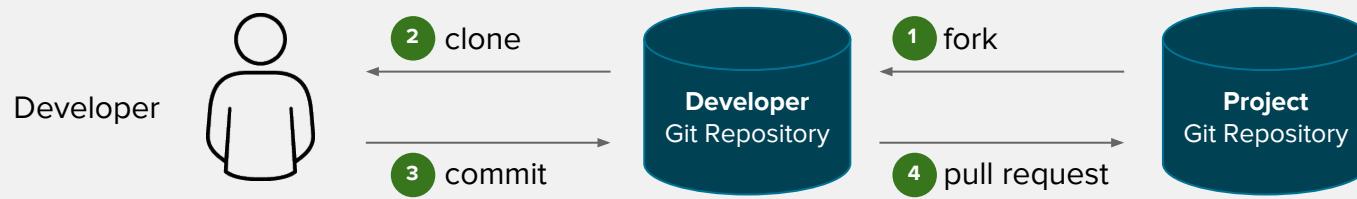
# COOLSTORE APPLICATION



The background of the slide features a dark teal overlay with a faint, abstract architectural pattern of overlapping building facades. The pattern consists of various geometric shapes and lines, creating a sense of depth and perspective.

# CREATING ON-DEMAND ENVIRONMENTS

# FORKING A GIT REPOSITORY

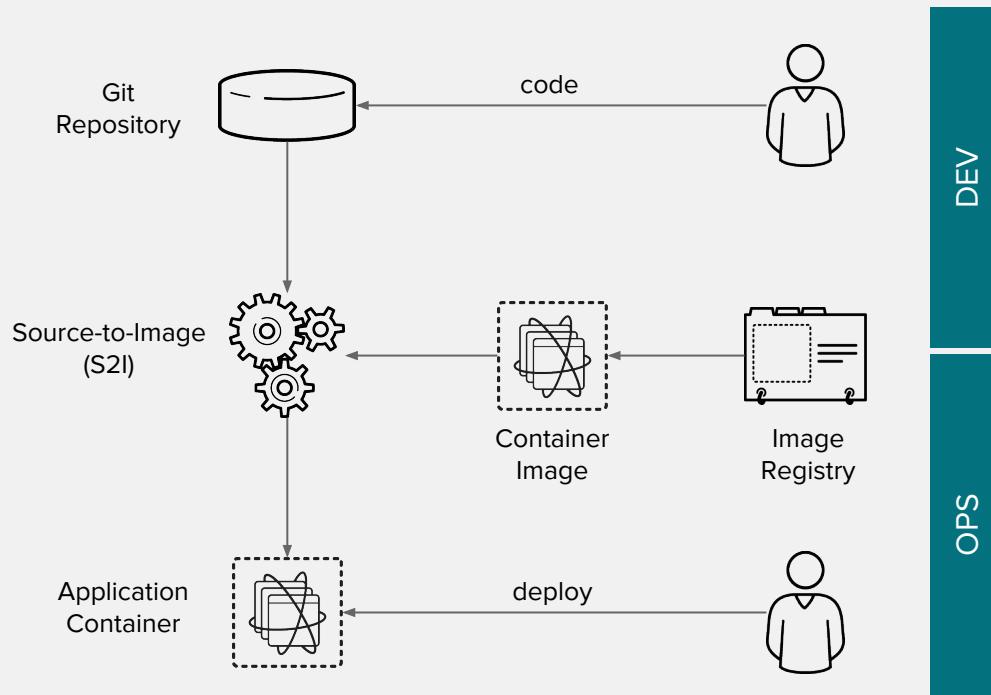


# BUILDING CONTAINER IMAGES WITH SOURCE-TO-IMAGE

**CODE**

**BUILD**

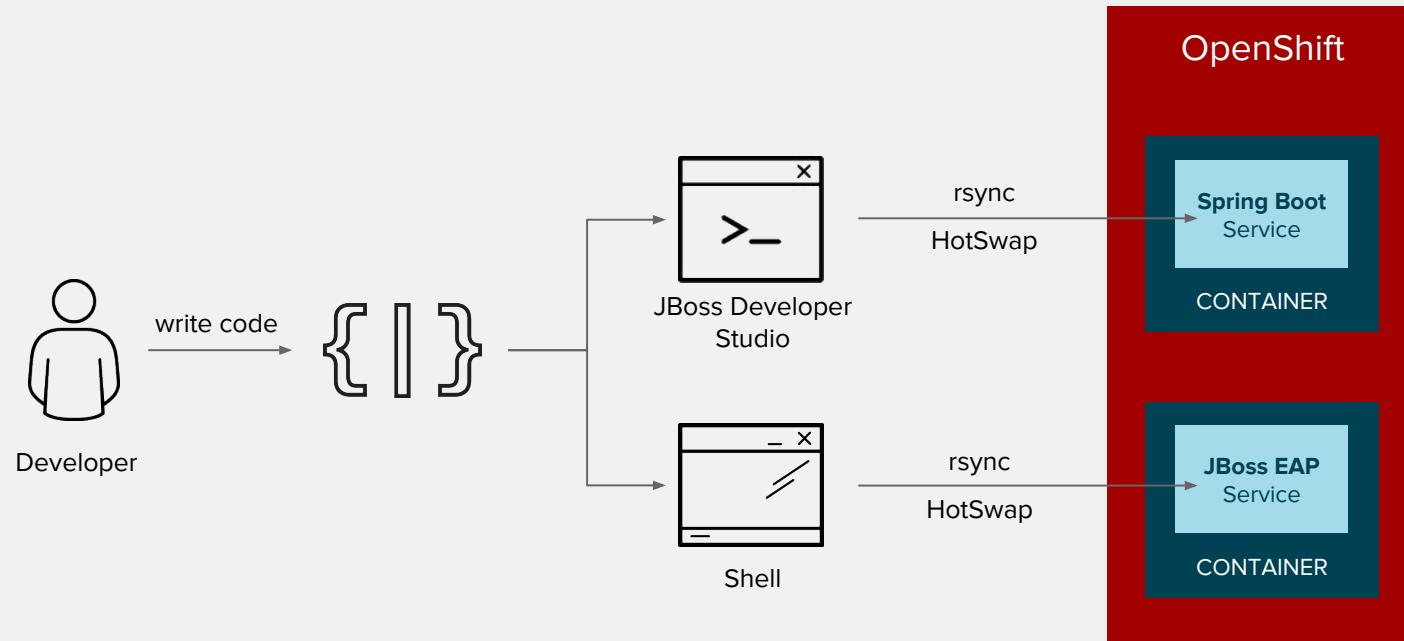
**DEPLOY**



The background of the slide features a dark teal overlay with a faint, stylized image of modern skyscrapers. The buildings are depicted with sharp, angular facades and windows, creating a sense of depth and urban architecture.

# DEVELOPER PRODUCTIVITY WITH CONTAINERS

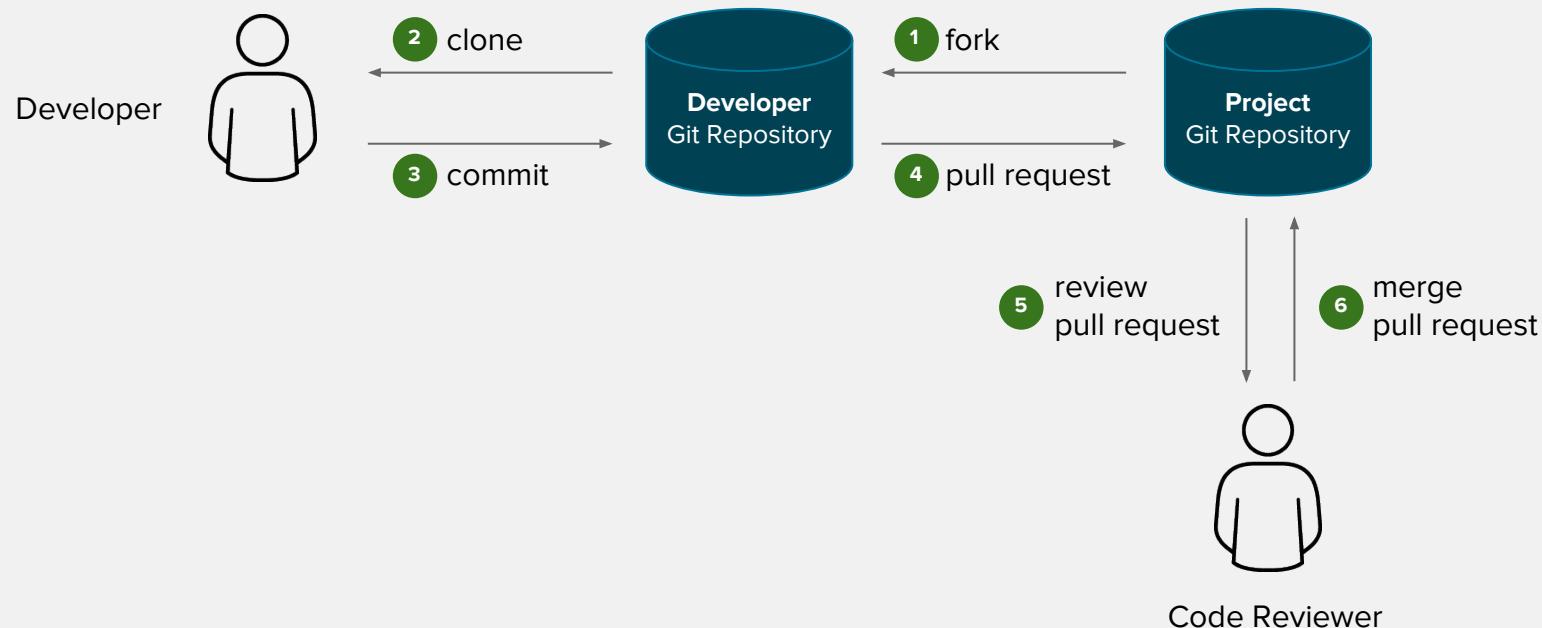
# DEPLOY CODE WITHOUT BUILD DELAYS



The background of the slide features a dark teal overlay with a faint, abstract architectural pattern of overlapping building facades. The pattern consists of various geometric shapes and lines, creating a sense of depth and perspective.

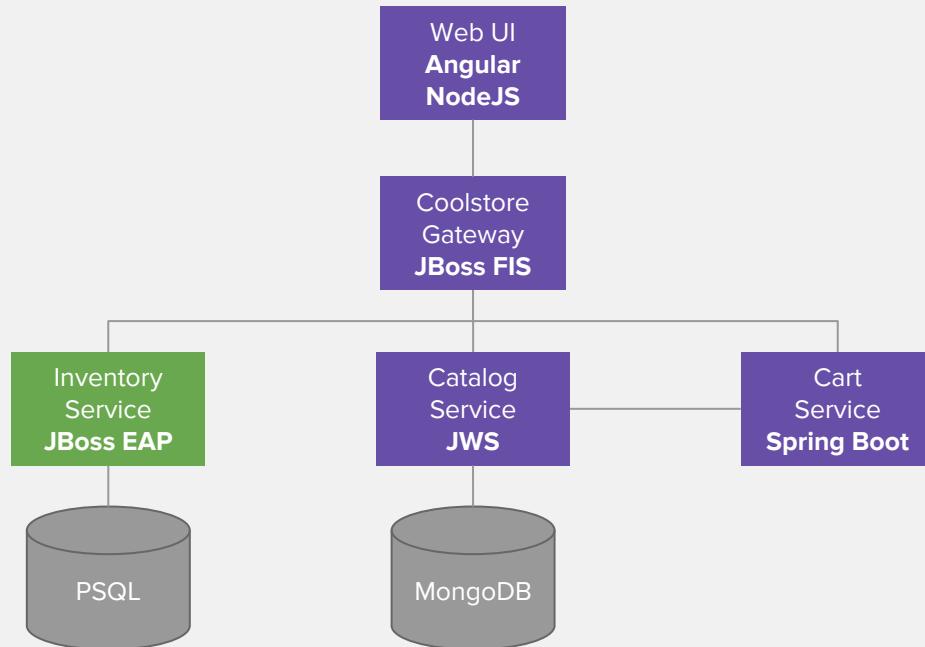
# BUILDING QUALITY INTO THE DEVELOPMENT PROCESS

# BUILDING QUALITY INTO THE PROCESS

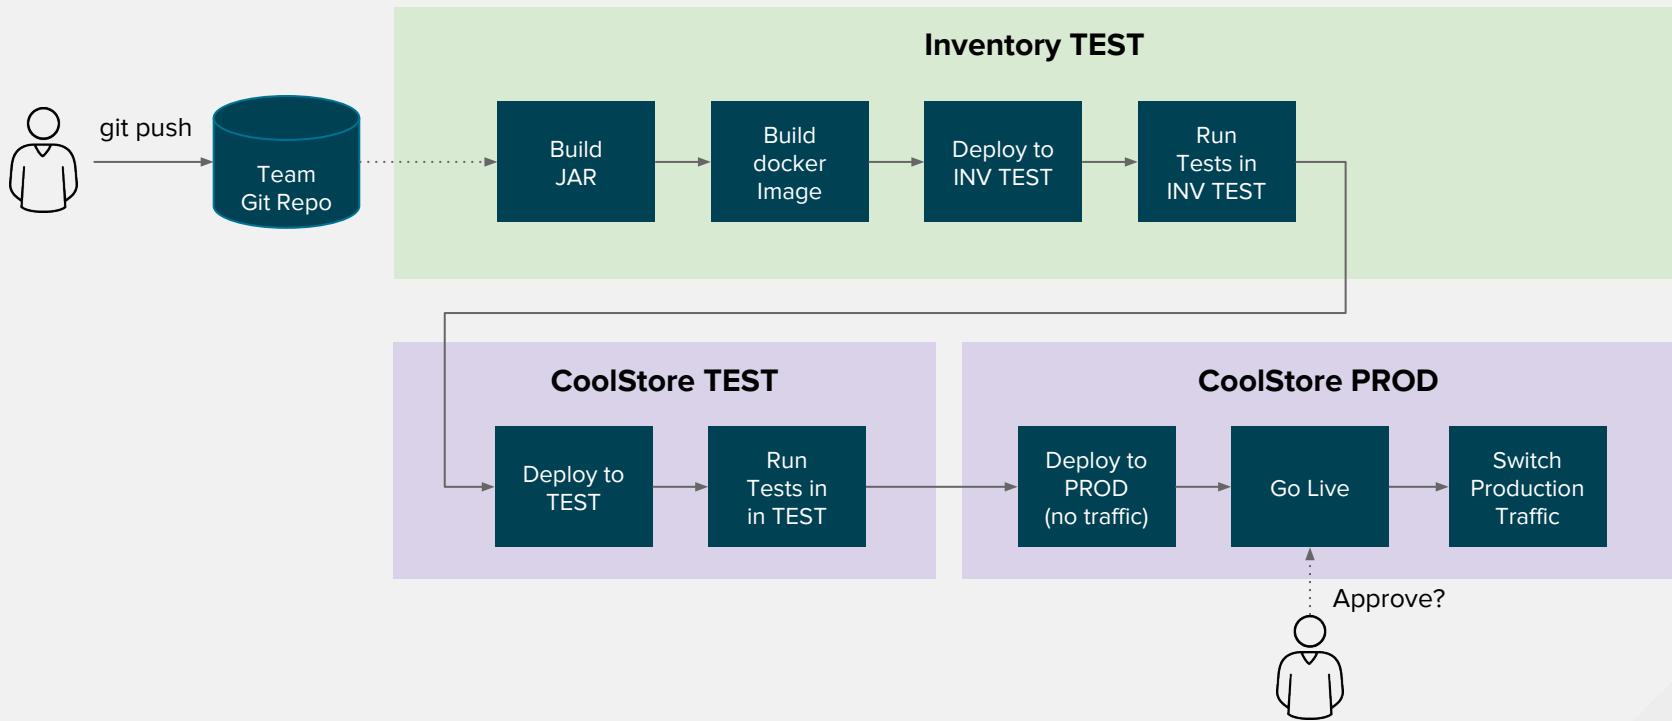


# **INCREASING DELIVERY SPEED WITH CONTINUOUS INTEGRATION AND CONTINUOUS DELIVERY**

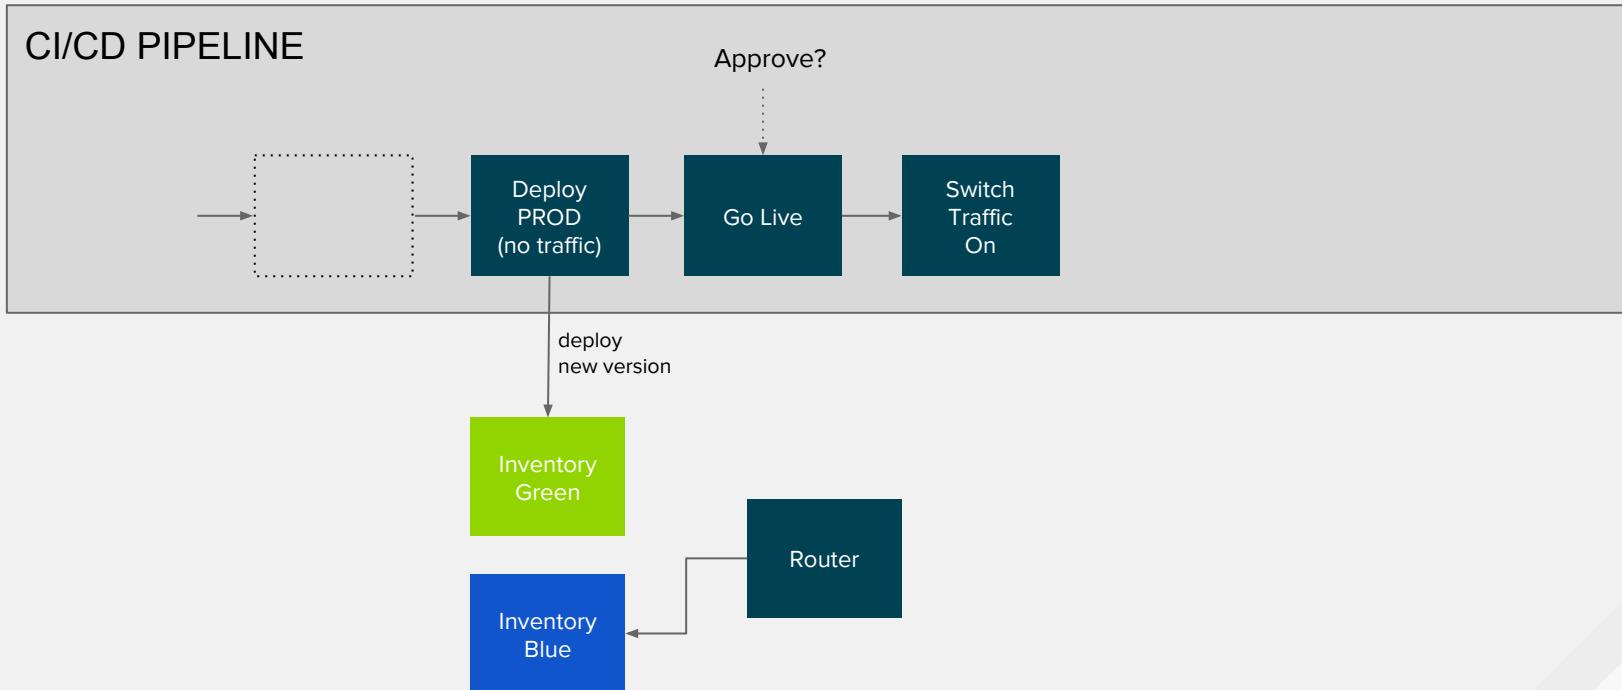
# COOLSTORE MICROSERVICES APPLICATION



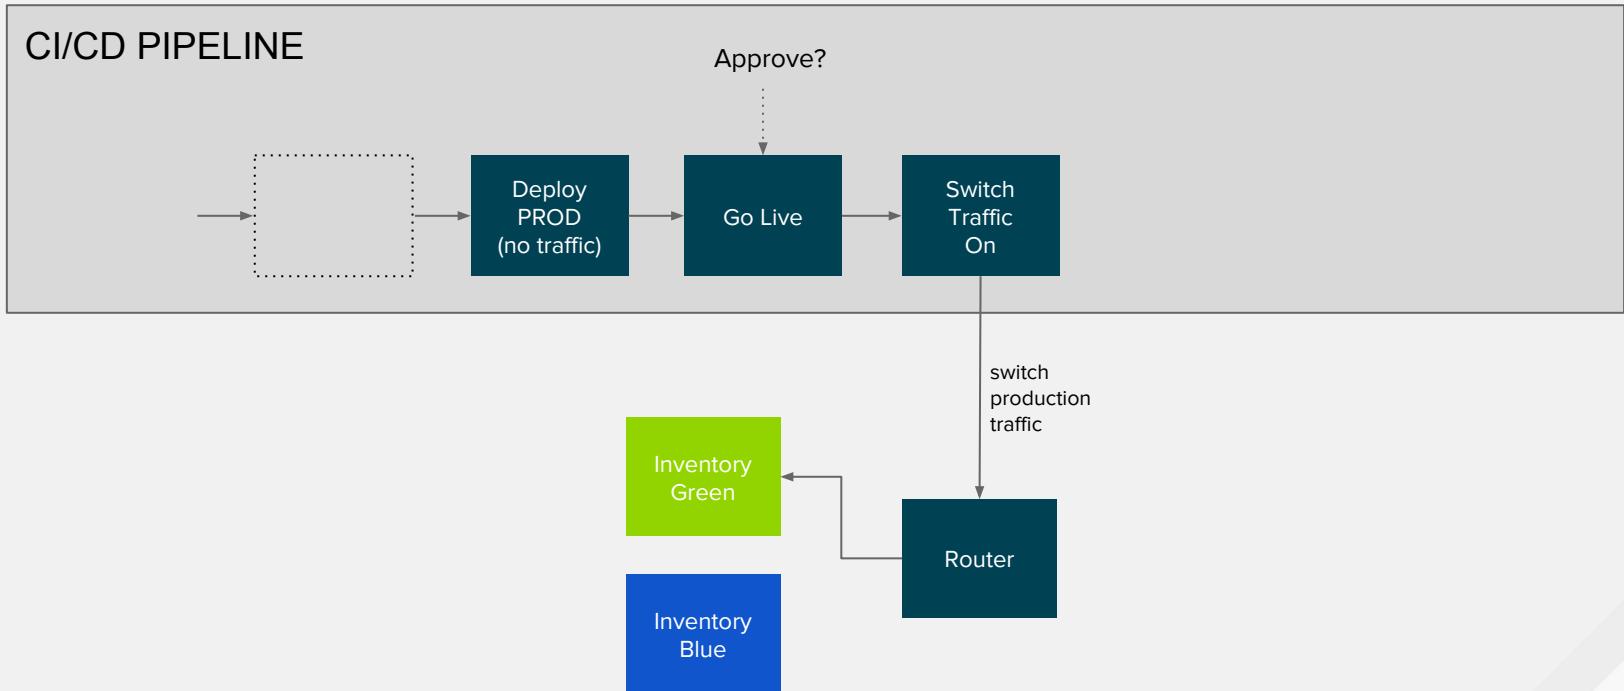
# CONTINUOUS DELIVERY PIPELINE



# ZERO DOWNTIME WITH BLUE/GREEN DEPLOYMENT



# ZERO DOWNTIME WITH BLUE/GREEN DEPLOYMENT





# **LAB 1**

# **BUILDING AND DEPLOYING A**

# **FAST-MOVING MONOLITH**

# FAST-MOVING MONOLITH

- Large organizations have a tremendous amount of resources **invested in existing** monolith applications
- Looking for a **sane way** to capture the benefits of containers and orchestration **without having to complete rewrite**
- **OpenShift** provides the platform for their existing investment with the benefit of a **path forward** for microservice based apps in the future

# **FAST-MOVING MONOLITH ADVANTAGES**

- Easier to develop since all dependencies are included
- Single code base for teams to work on
- No API backwards compatibility issues since all logic is packaged with the application
- Single deployable unit

# FAST-MOVING MONOLITH DISADVANTAGES

- Hard for large teams to all work in the same code base. *Who broke the build?*
- Longer release cycles and even small changes goes through entire test suite and validation
- Same language / Same Framework in most cases
- *Entropic resistance* is low

# LAB 1 - FAST-MOVING MONOLITH

- In this lab, the coolstore monolith will be built and deployed to OpenShift from your local workstation demonstrating a typical Java application developer workflow
- A sample pipeline is included which will be used to deploy across dev and prod environment
- First, deploy the coolstore monolith dev project (don't forget to include `-b app-partner` when you run `git clone` - this is the branch in use for this lab!):

```
$ mkdir ~/coolstore
$ cd ~/coolstore
$ git clone -b app-partner https://github.com/modernize-legacy-apps/monolith
$ cd monolith
$ oc new-project coolstore-<USERNAME>
$ oc process -f src/main/openshift/template.json | oc create -f -
```

# LAB 1 - FAST-MOVING MONOLITH

- Notice there is **no deployment yet** as there is no build yet:

Project  
coolstore

Overview

Applications >

Builds >

Resources >

Storage

No deployments.

A new deployment will start automatically when an image is available for [coolstore/coolstore:latest](#).

coolstore-dev

Deployment Config [coolstore-dev](#)

coolstore-dev-postgresql

Deployment Config [coolstore-dev-postgresql](#) - 2 days ago

CONTAINER: COOLSTORE-DEV-POSTGRESQL

Image: centos/postgresql-95-centos7

Ports: 5432/TCP

1 pod

# LAB 1 - FAST-MOVING MONOLITH

- A number of OpenShift objects were created:
  - A Postgres database deployment, service, and route
  - A coolstore *binary* BuildConfig
    - Binary builds accept source code through stdin in later commands (vs. kicking off a build within an OpenShift pod)
  - Services, Routes and DeploymentConfig for coolstore-dev
  - Service Accounts and Secrets (for cluster permissions)
- Build the coolstore monolith (.war file) locally and deploy into the OpenShift via the binary build:

```
$ cd ~/coolstore/monolith
$ mvn clean package -Popenshift
$ oc start-build coolstore --from-file deployments/ROOT.war --follow
```

# LAB 1 - FAST-MOVING MONOLITH

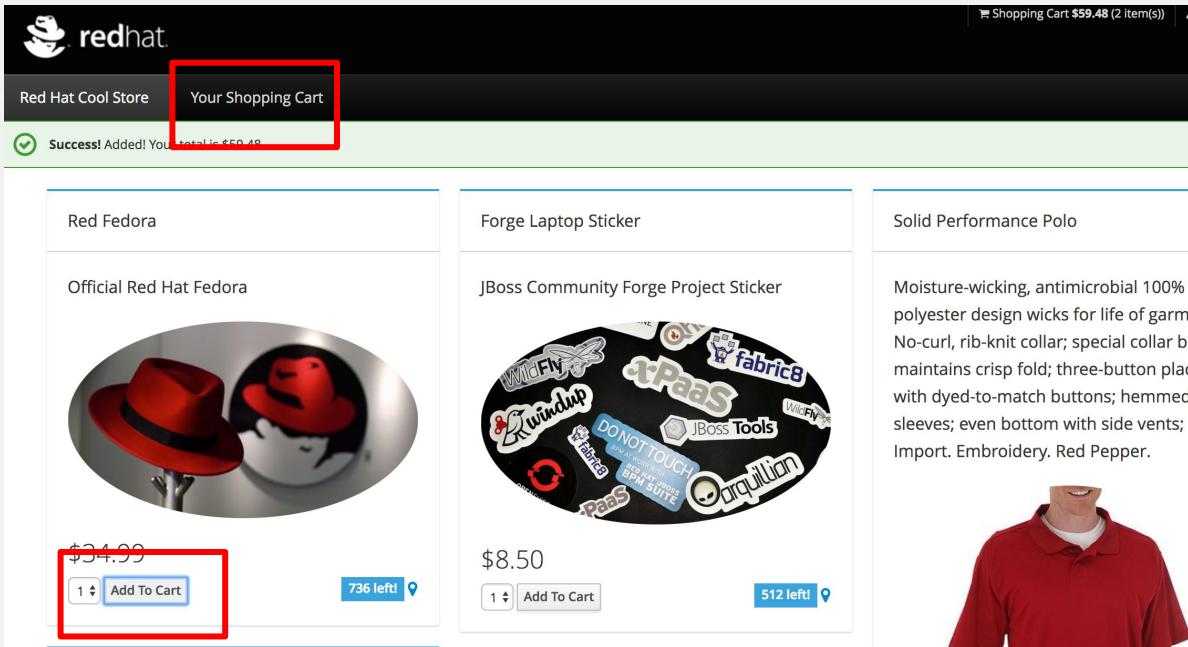
- Binary builds still produce Linux container images and are stored into the container registry.
- In the web console, navigate to *Builds* → *Images* to see the new image created as a result of combining the monolith .war file with the JBoss EAP builder image:

The screenshot shows the Red Hat OpenShift web console interface. On the left, there is a sidebar with navigation links: Overview, Applications, Builds (which is currently selected), Resources, and Storage. The main content area is titled "Image Streams" and displays a table of Docker images. The table has columns: Name, Docker Repo, Tags, and Updated. One row in the table is highlighted with a red border. The highlighted row contains the following information: Name: coolstore, Docker Repo: 172.30.147.123:5000/coolstore/coolstore, Tags: latest, and Updated: 2 days ago.

Name	Docker Repo	Tags	Updated
coolstore	172.30.147.123:5000/coolstore/coolstore	latest	2 days ago

# LAB 1 - FAST-MOVING MONOLITH

- Once the deployment is complete, navigate to the application by clicking on its route in the OpenShift web console Overview and exercise the app by adding/removing products to your shopping cart:



# LAB 1 - FAST-MOVING MONOLITH

- The app contains an AngularJS web frontend which makes REST calls to its backend (e.g. /services/products)
- The app's backend contains services implement as Stateless and Stateful EJBs backed by JPA entities
- The product catalog and inventory are stored together in the database.
- Log into the database pod (with postgresql in the pod name) and inspect the values stored (you will need to copy/paste the password from the POSTGRESQL\_PASSWORD environment variable when running the psql command):

```
% oc env dc/coolstore-dev-postgresql --list      # copy value of POSTGRESQL_PASSWORD to clipboard
$ oc get pods
$ oc rsh coolstore-dev-postgresql-1-3rfuw
sh-4.2$ psql -h $HOSTNAME -d $POSTGRESQL_DATABASE -U $POSTGRESQL_USER
Password for user userV31:XXXXXXXX
psql (9.5.4)
Type "help" for help.

monolith=> select * from INVENTORY;
monolith=> select * from PRODUCT_CATALOG;
```

# LAB 1 - FAST-MOVING MONOLITH

- The `coolstore-dev` build and deployment is responsible for building new versions of the app after code changes and deploying them to the development environment.
- Production environments will ideally use different projects, and even different nodes in an OpenShift cluster. For simplicity we have dev and prod in the same project.
- Create the production deployment objects that will be used to promote builds from dev to prod using the supplied CI/CD pipeline:

```
$ oc process -f ~/coolstore/monolith/src/main/openshift/template-prod.json | oc create -f -
```

# LAB 1 - FAST-MOVING MONOLITH

- Several types of OpenShift objects are created:
  - A production Postgres database deployment, service, and route
  - Services, Routes and DeploymentConfig for coolstore-prod (production env)
  - Service Accounts and Secrets (for cluster permissions)
  - A CI/CD server (Jenkins)
  - A CI/CD pipeline
- Notice that **no deployment occurs** for the production version of the app.
- Deployment of the production app occurs as a result of images from the dev environment being appropriately tagged. Inspect the pipeline and observe the steps it executes during the pipeline run. Notice at the end of the script, the latest coolstore image is tagged with :prod, triggering a production deployment:

```
$ oc get bc/monolith-pipeline -o yaml
...
    openshiftTag(sourceStream: 'coolstore', sourceTag: 'latest', namespace: '', destinationStream: 'coolstore',
destinationTag: ''prod'', destinationNamespace: '')
...
...
```

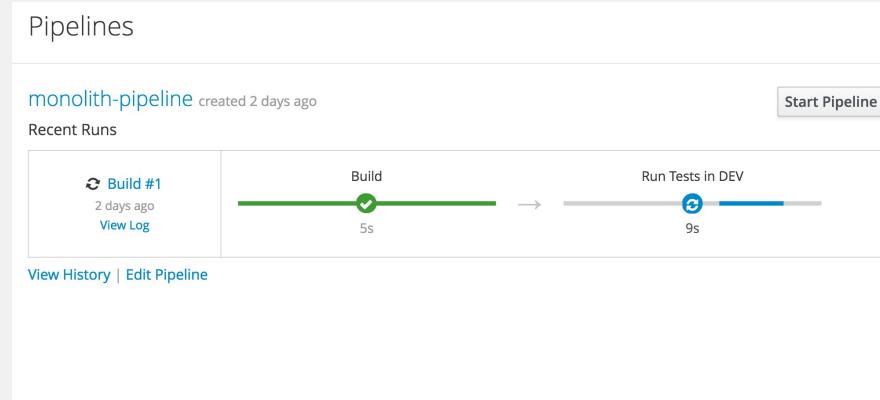
# LAB 1 - FAST-MOVING MONOLITH

- Wait for the `jenkins` service to be completely available, then execute the pipeline by navigating to *Builds*→*Pipelines* and click on *Start Pipeline* next to the Monolith Pipeline:

The screenshot shows the Jenkins interface for the 'coolstore' project. The top navigation bar includes a home icon, the project name 'coolstore', a dropdown for 'Add to project', and user information for 'admin'. On the left, a sidebar lists 'Overview', 'Applications >', 'Builds >' (which is highlighted in blue), and 'Resources >'. The main content area is titled 'Pipelines' and displays a single pipeline named 'monolith-pipeline', which was created 2 days ago. Below the pipeline name, it says 'No pipeline builds have run for monolith-pipeline.' To the right of the pipeline name is a 'Start Pipeline' button, which is enclosed in a red rectangular box.

# LAB 1 - FAST-MOVING MONOLITH

- Observe the pipeline as it executes its stages
- Once the pipeline is complete, the production app is deployed.
- Typically there would be a human approval step before going live
- Pipelines are also good at executing advanced deployment strategies (blue/green, canary)
- Return to the Overview page, wait for the deployment to complete, then click on the route for the `coolstore-prod` deployment to test the production app (identical to `coolstore-dev`).



# LAB 1 - FAST-MOVING MONOLITH

- You should now have two identical copies of the app deployed, along with two databases:

The screenshot shows the OpenShift web console interface for the 'coolstore' project. The left sidebar includes links for Home, Overview, Applications, Builds, Resources, Storage, and Monitoring. The main area displays two identical deployment configurations, each consisting of a deployment config and a PostgreSQL database deployment config.

**Dev Environment:**

- Deployment Config coolstore-dev:** CONTAINER: COOLSTORE-DEV. Image: coolstore/coolstore. Ports: 8080/TCP (http), 8443/TCP (https), 8778/TCP (loki), 8888/TCP (ping). Status: Complete, 2 days ago.
- Deployment Config coolstore-dev-postgresql:** CONTAINER: COOLSTORE-DEV-POSTGRESQL. Image: centos/postgresql-95-centos7. Ports: 5432/TCP. Status: Complete, 2 days ago.

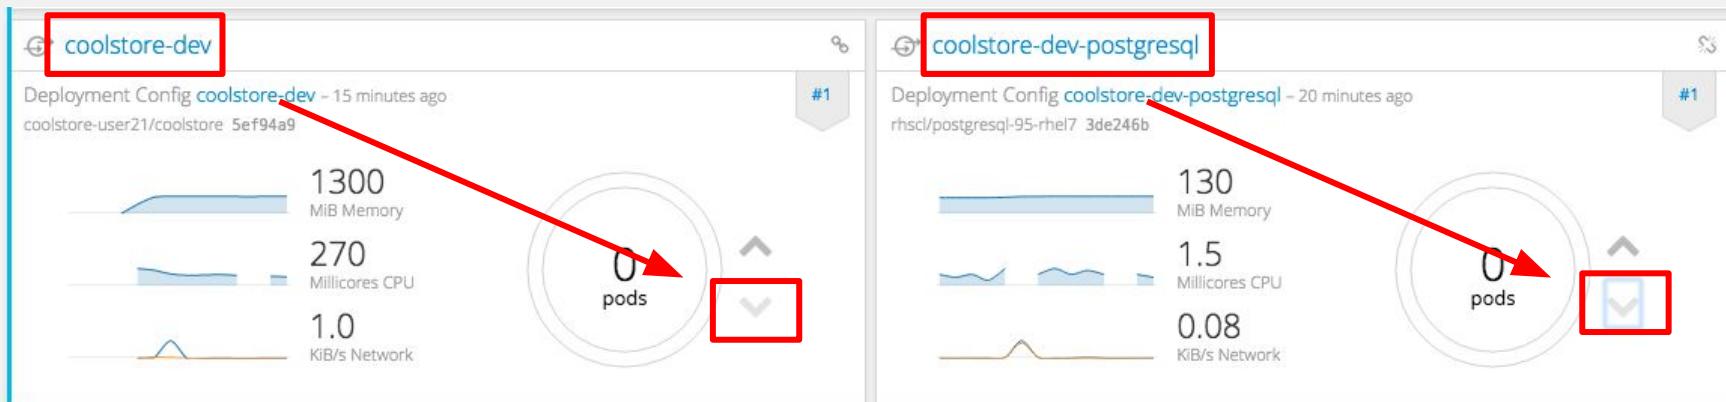
**Production Environment:**

- Deployment Config coolstore-prod:** CONTAINER: COOLSTORE-PROD. Image: coolstore/coolstore. Ports: 8080/TCP (http), 8443/TCP (https), 8778/TCP (loki), 8888/TCP (ping). Status: Complete, 2 days ago.
- Deployment Config coolstore-prod-postgresql:** CONTAINER: COOLSTORE-PROD-POSTGRESQL. Image: centos/postgresql-95-centos7. Ports: 5432/TCP. Status: Complete, 2 days ago.

A green box highlights the Dev environment, and a blue box highlights the Production environment. Labels '← Dev' and '← Prod' are positioned to the right of their respective boxes.

# LAB 1 - FAST-MOVING MONOLITH

- (Optional) If your OpenShift environment seems slow, you can scale down the coolstore-dev environment by clicking the down arrow next to the coolstore-dev and coolstore-dev-postgresql pods:



- (Optional) If your the build pipeline fails to start, you can manually tag the images:

```
$ oc tag coolstore:latest coolstore:prod
```

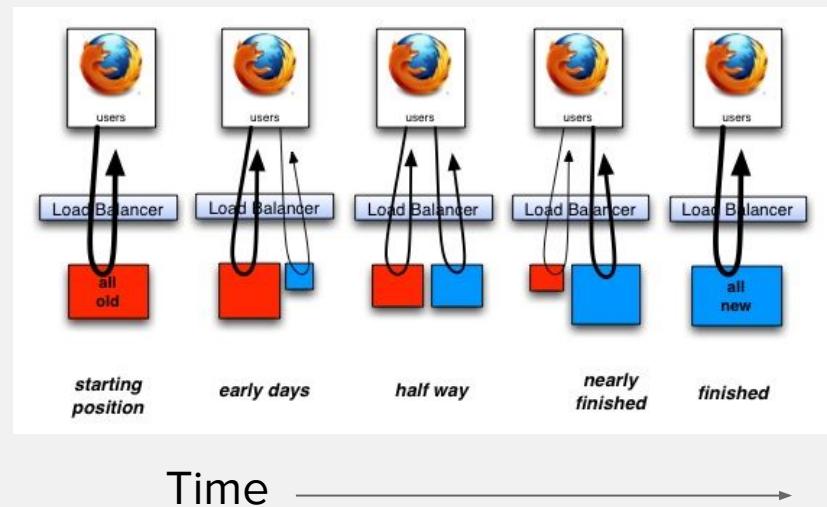


# **LAB 2**

# **STRANGLING THE MONOLITH**

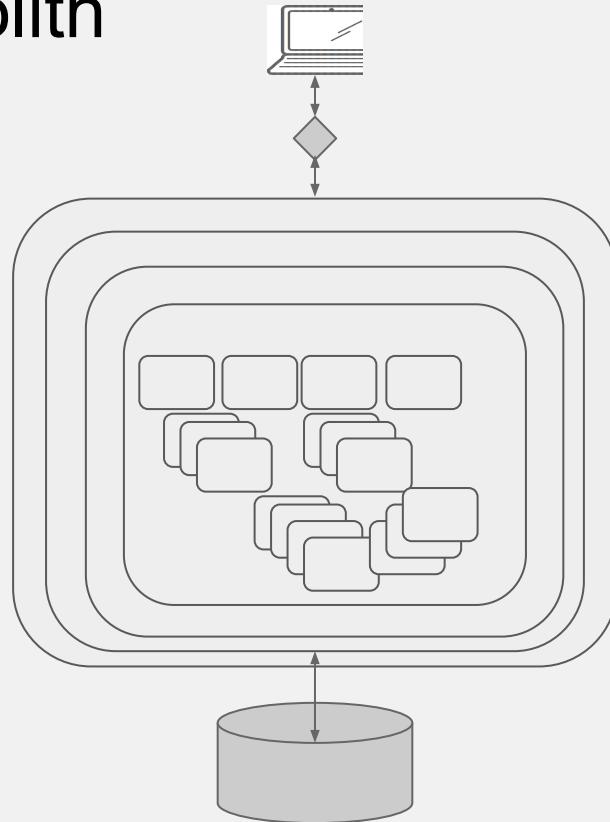
# LAB 2 - STRANGLING THE MONOLITH

- Strangling - **incrementally** replacing functionality in app with something better (cheaper, faster, easier to maintain).
- As functionality is replaced, “dead” parts of monolith can be removed/retired.
- You can also wait for all functionality to be replaced before retiring anything!
- You can optionally include new functionality during strangulation to make it more attractive to business stakeholders.

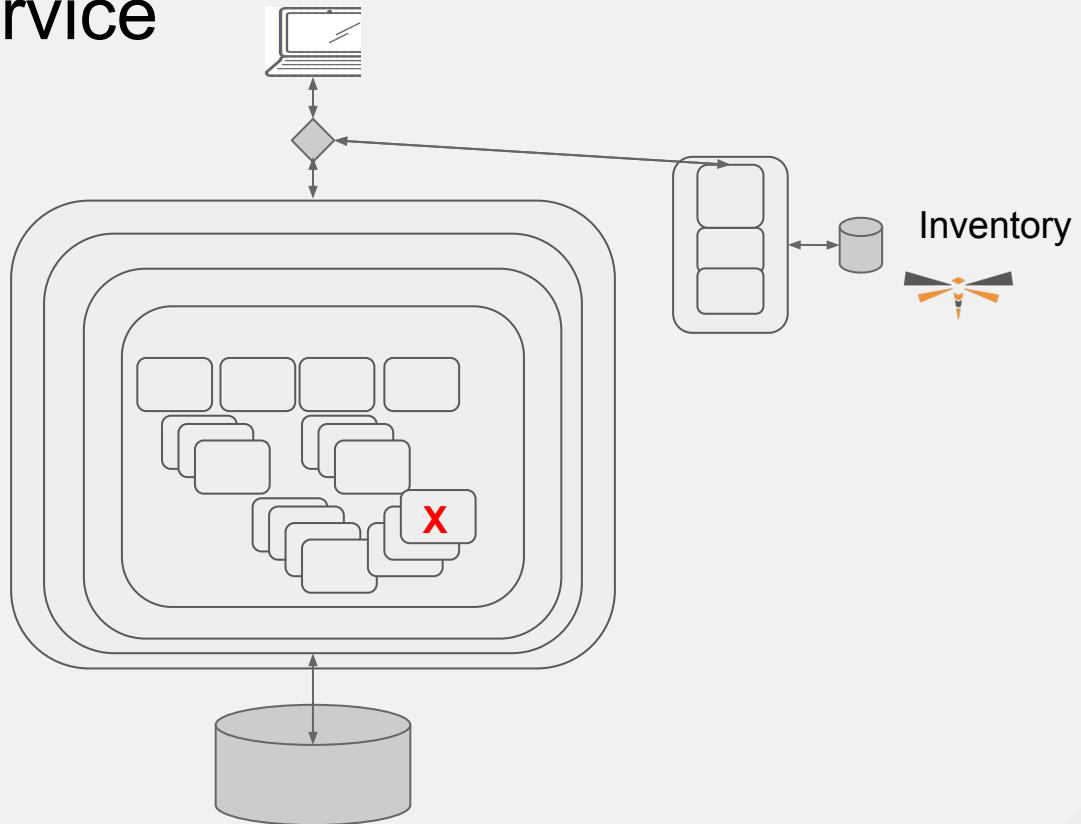


<http://paulhammant.com/2013/07/14/legacy-application-strangulation-case-studies>

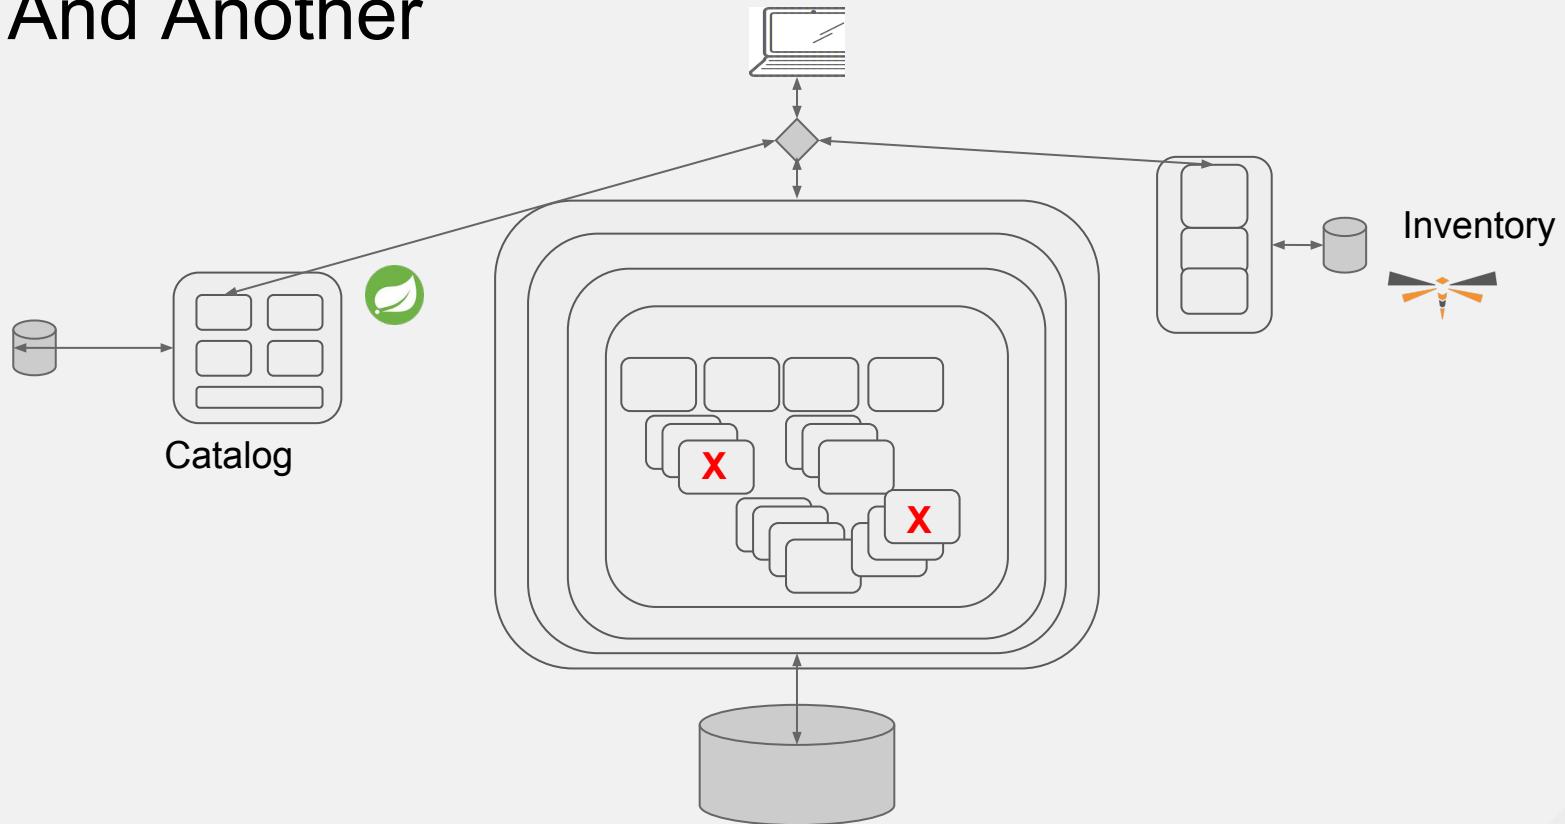
# 1) Strangle Monolith



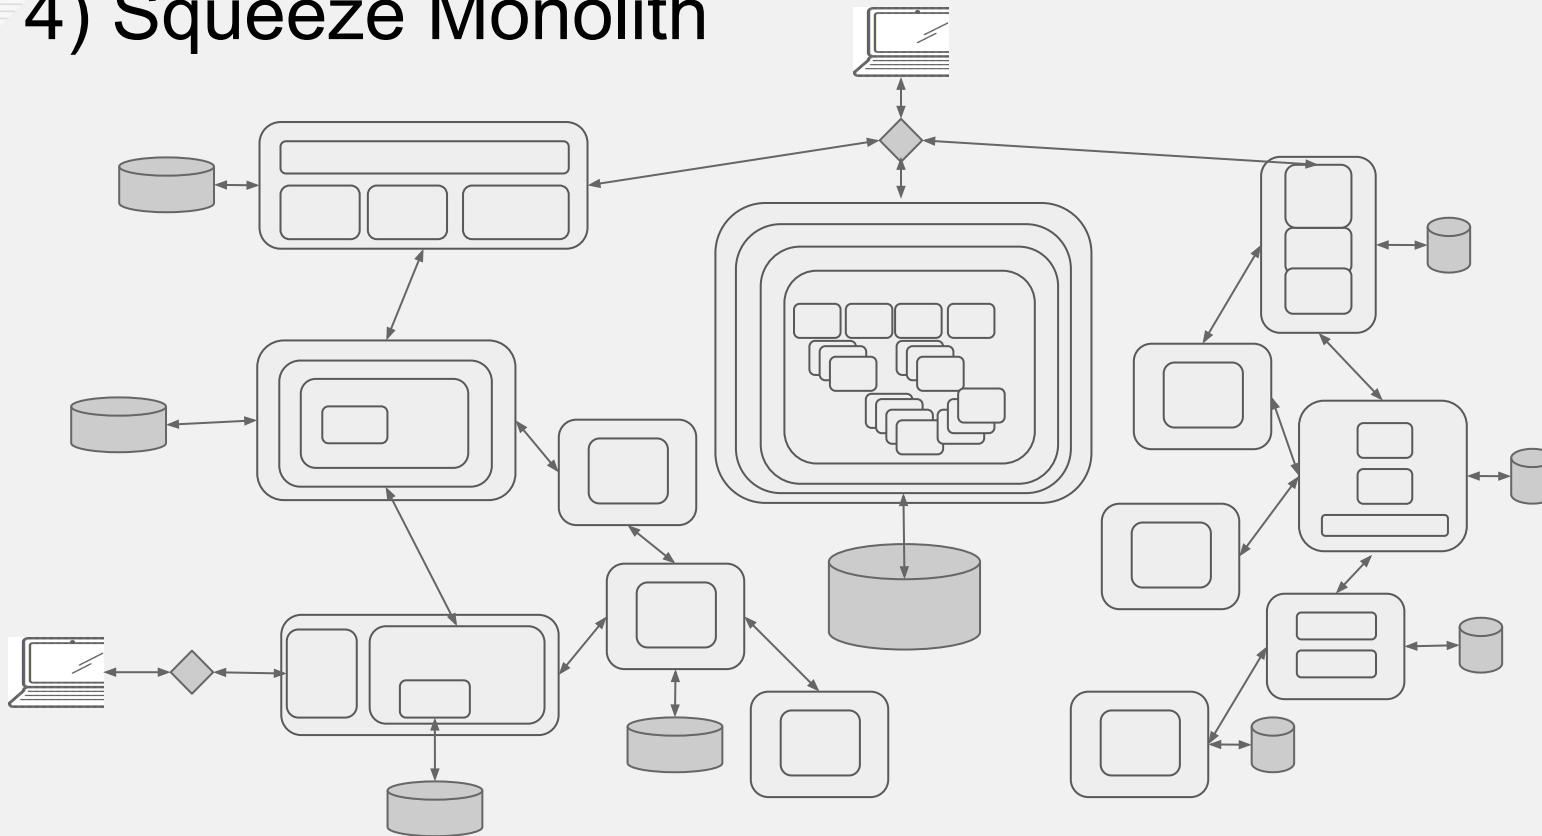
## 2) Add a Microservice



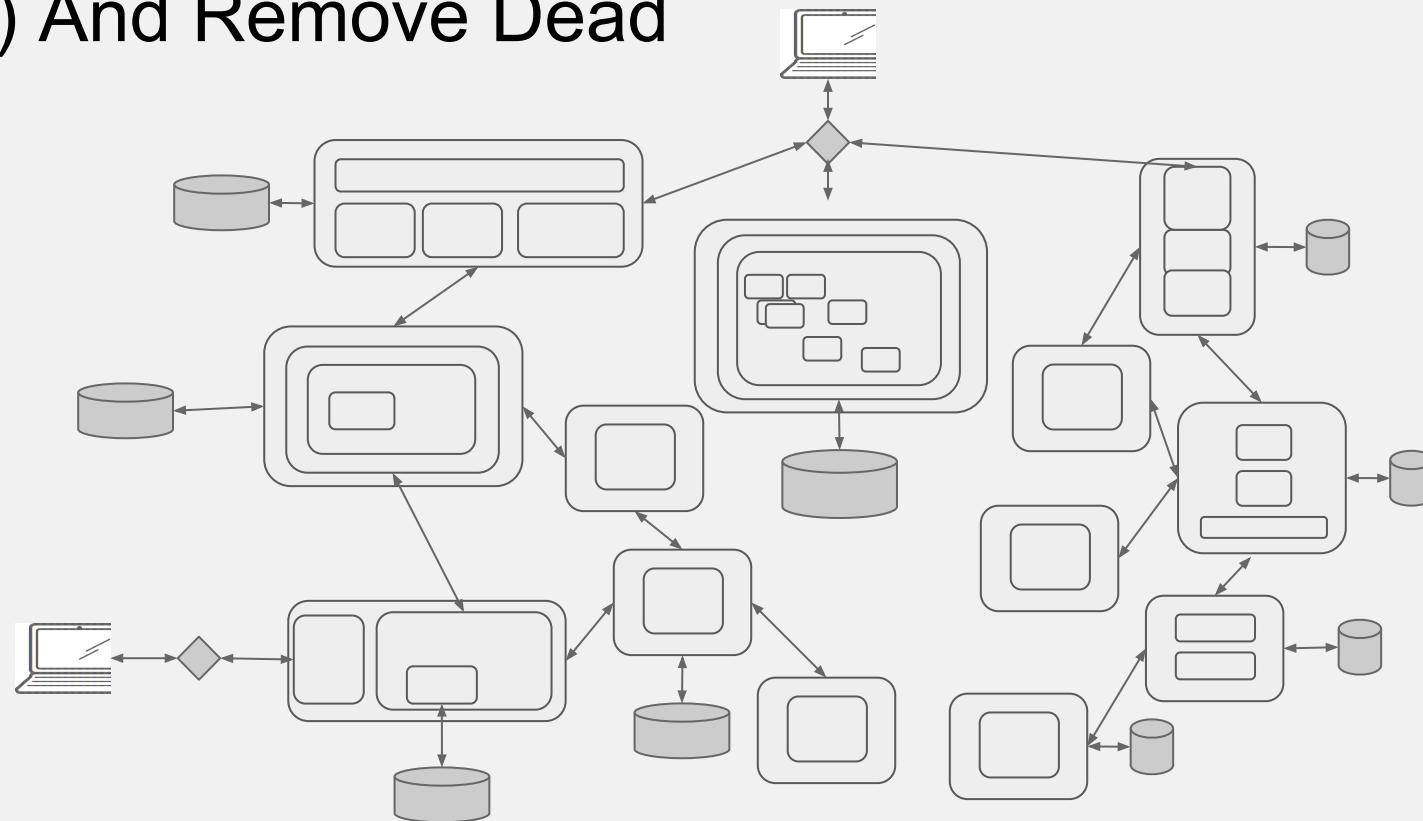
# 3) And Another



# 4) Squeeze Monolith



# 5) And Remove Dead



# LAB 2 - STRANGLING THE MONOLITH

- In this lab, you will begin to ‘strangle’ the coolstore monolith by implementing two of its services as external microservices, split along business boundaries
- Once implemented, traffic destined to the original monolith’s services will be redirected (via OpenShift software-defined routing) to the new services
- First, build and deploy the catalog microservice (Based on Spring Boot)
- We will be using the Fabric8 Maven Plugin to automatically deploy to openshift

```
$ cd ~/coolstore
$ git clone -b app-partner https://github.com/modernize-legacy-apps/catalog
$ cd catalog
$ oc project coolstore-<USERNAME>
$ mvn clean fabric8:deploy -Popenshift -DskipTests
```

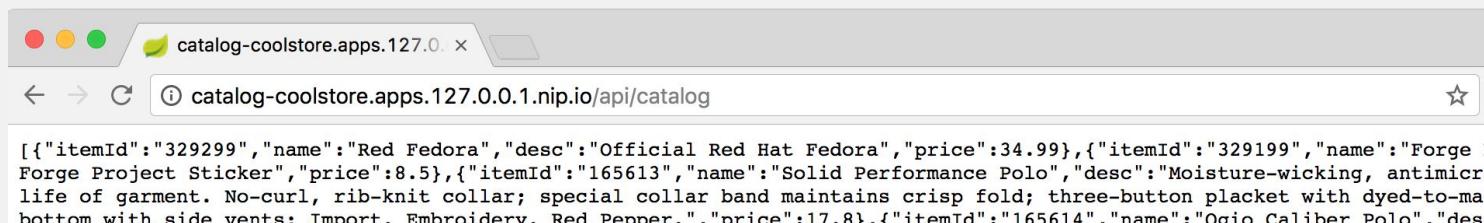
# LAB 2 - STRANGLING THE MONOLITH

- Once the new catalog service is deployed, test it with curl:

```
$ oc get route/catalog
NAME      HOST/PORT          PATH      SERVICES  PORT      TERMINATION
catalog   catalog-coolstore.apps.127.0.0.1.nip.io      catalog    8080

$ curl http://catalog-coolstore.apps.127.0.0.1.nip.io/api/catalog
[{"itemId": "329299", "name": "Red Fedora", "desc": "Official Red Hat
Fedora", "price": 34.99}, {"itemId": "329199", "name": "Forge Laptop Sticker", "desc": "JBoss Community Forge Project
Sticker", "price": 8.5}
...
```

- You can also test it by clicking on the route to the catalog service and adding /api/catalog to the URL:



# LAB 2 - STRANGLING THE MONOLITH

- Now we will do the same for the Inventory service, based on WildFly Swarm:

```
$ cd ~/coolstore
$ git clone -b app-partner https://github.com/modernize-legacy-apps/inventory
$ cd inventory
$ oc project coolstore-<USERNAME>
$ mvn clean fabric8:deploy -Popenshift -DskipTests
```

- The catalog microservice is responsible for retrieving and returning a list of products and their descriptions (price, images, etc).
- The inventory service is responsible for retrieving and returning inventory for a given product.
- The catalog service does not specifically know about or care about inventory. In the next lab we will combine the two!

# LAB 2 - STRANGLING THE MONOLITH

- The inventory microservice has a separate database from the current monolith.
- This allows teams split along business concerns to independently develop, test, deploy and scale the service and its database
- Once the new inventory service is deployed, dump its database:

```
% oc env dc/inventory-database --list # get the POSTGRESQL_PASSWORD
$ oc get pods
$ oc rsh inventory-database-1-kkxs2
sh-4.2$ psql -h $HOSTNAME -d $POSTGRESQL_DATABASE -U $POSTGRESQL_USER
Password for user userV31:XXXXXXXX
psql (9.5.4)
Type "help" for help.

monolith=> select * from inventory;
      itemid |           link           | location | quantity
-----+-----+-----+-----+
  329299 | http://maps.google.com/?q=Raleigh | Raleigh |      736
...
...
```

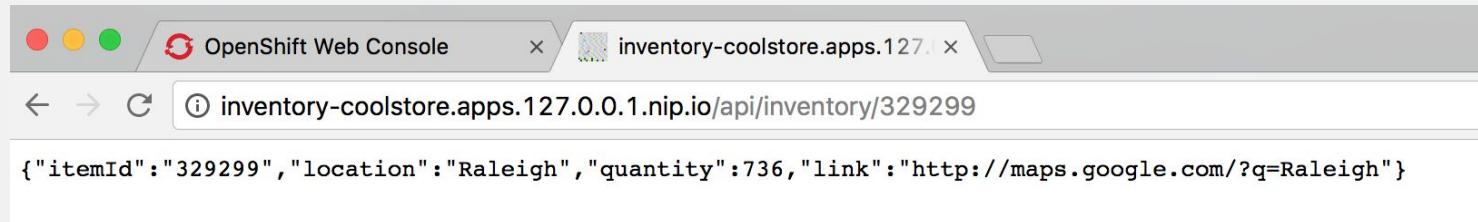
# LAB 2 - STRANGLING THE MONOLITH

- Once the new service is deployed, test it with curl
- You should get a JSON object representing the inventory availability of product 329299

```
$ oc get route/inventory
NAME      HOST/PORT          PATH      SERVICES    PORT      TERMINATION
inventory  inventory-coolstore.apps.127.0.0.1.nip.io      inventory  8080

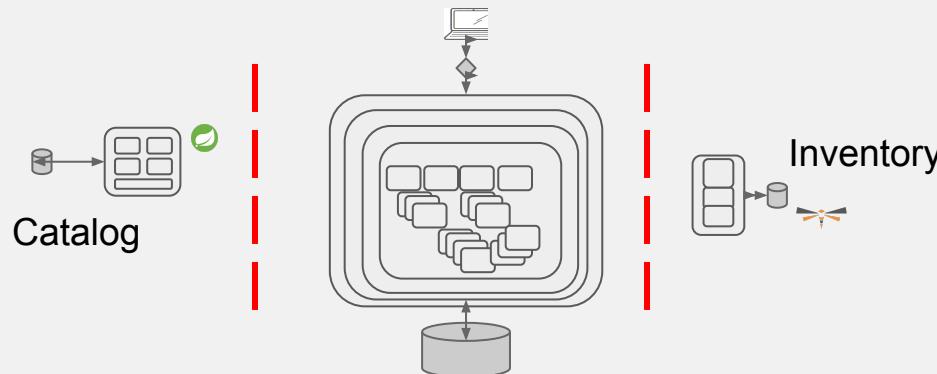
% curl http://inventory-coolstore.apps.127.0.0.1.nip.io/api/inventory/329299
{"itemId":"329299","location":"Raleigh","quantity":736,"link":"http://maps.google.com/?q=Raleigh"}
```

- You can also test it by clicking on the route to the inventory service and adding /api/inventory/329299 to the URL:



# LAB 2 - STRANGLING THE MONOLITH

- You have begun to replace your monolith's services with new, lightweight Java-based microservices - congratulations!
- Although deployed, no web traffic will ever hit these new services as nothing (including the existing UI) knows of their existence.
- Microservice architectures have much more deployment flexibility but at the cost of increased complexity as the system becomes more distributed.
- In the next lab we'll discuss how to integrate the new microservices into the existing app while the development teams maintain complete development and deployment autonomy.



The background of the slide features a dark teal overlay with a faint, abstract architectural pattern of overlapping building facades and windows.

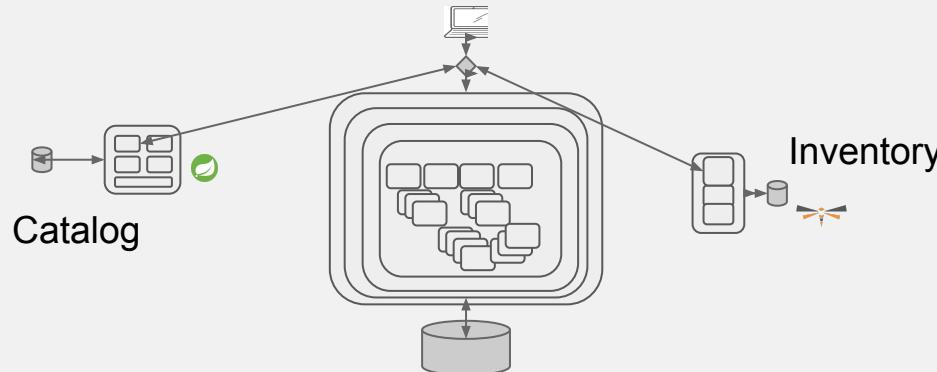
# LAB 3

# MICROSERVICE INTEGRATION

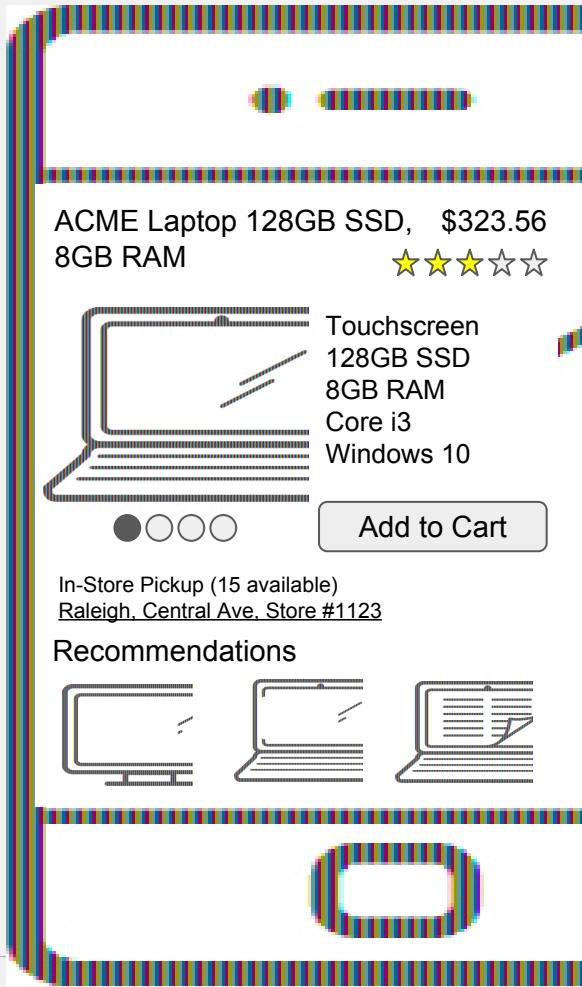
# PATTERNS

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- In previous labs, we created two new microservices with the intention of replacing functionality (“strangling”) the monolithic application.
- Currently no traffic is routed to them.
- If you were to re-route traffic from the monolith’s /services/products API to the new catalog service’s /services/catalog endpoint, you would be missing the inventory data.
- In this lab we will consider different options and architectures for integrating the microservices’ functionality into our app.

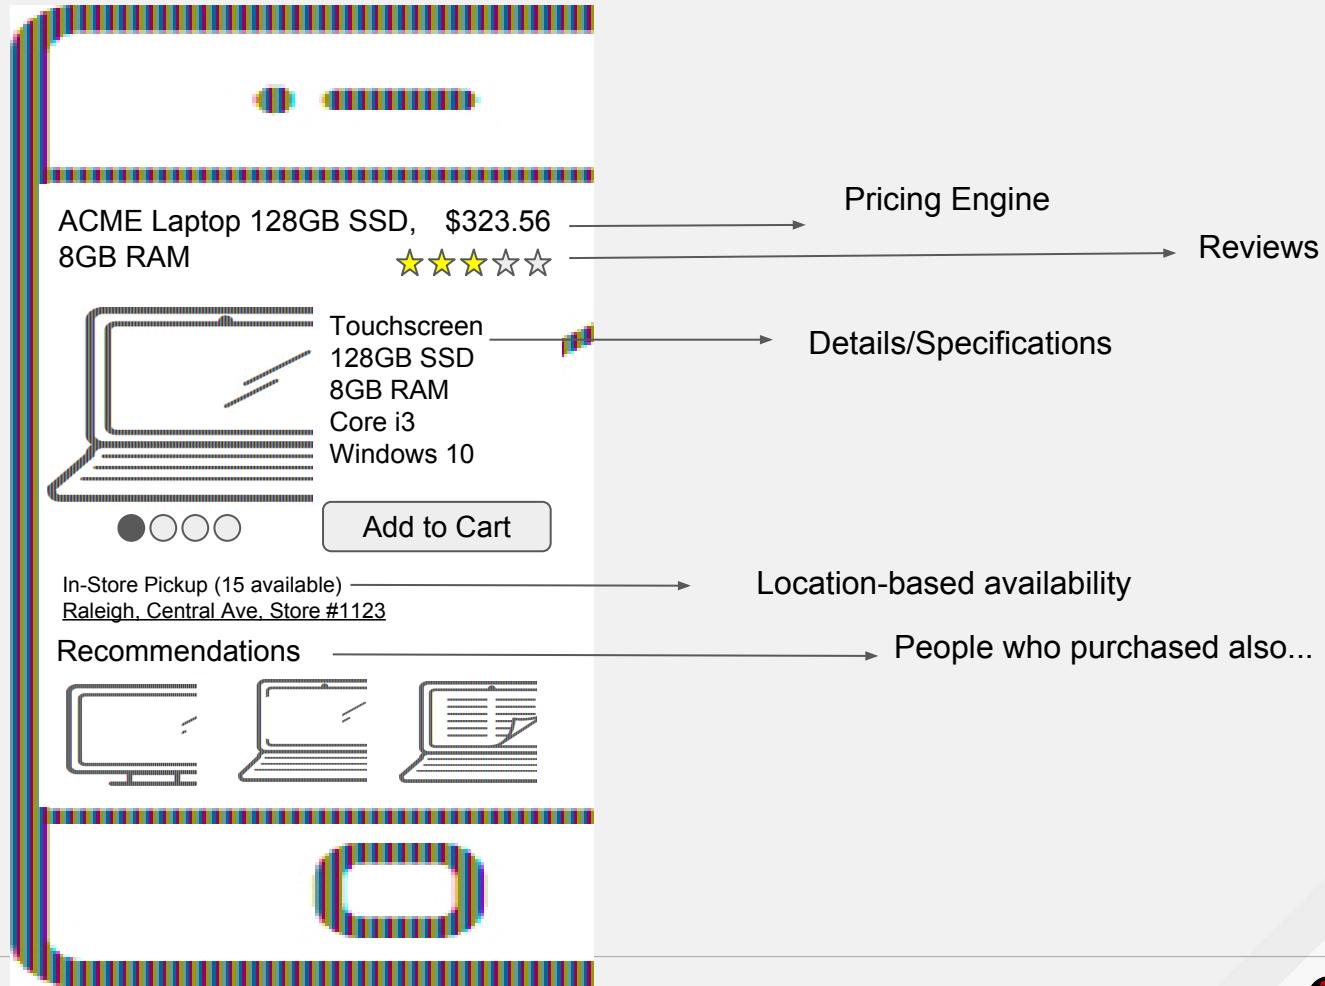


# Example

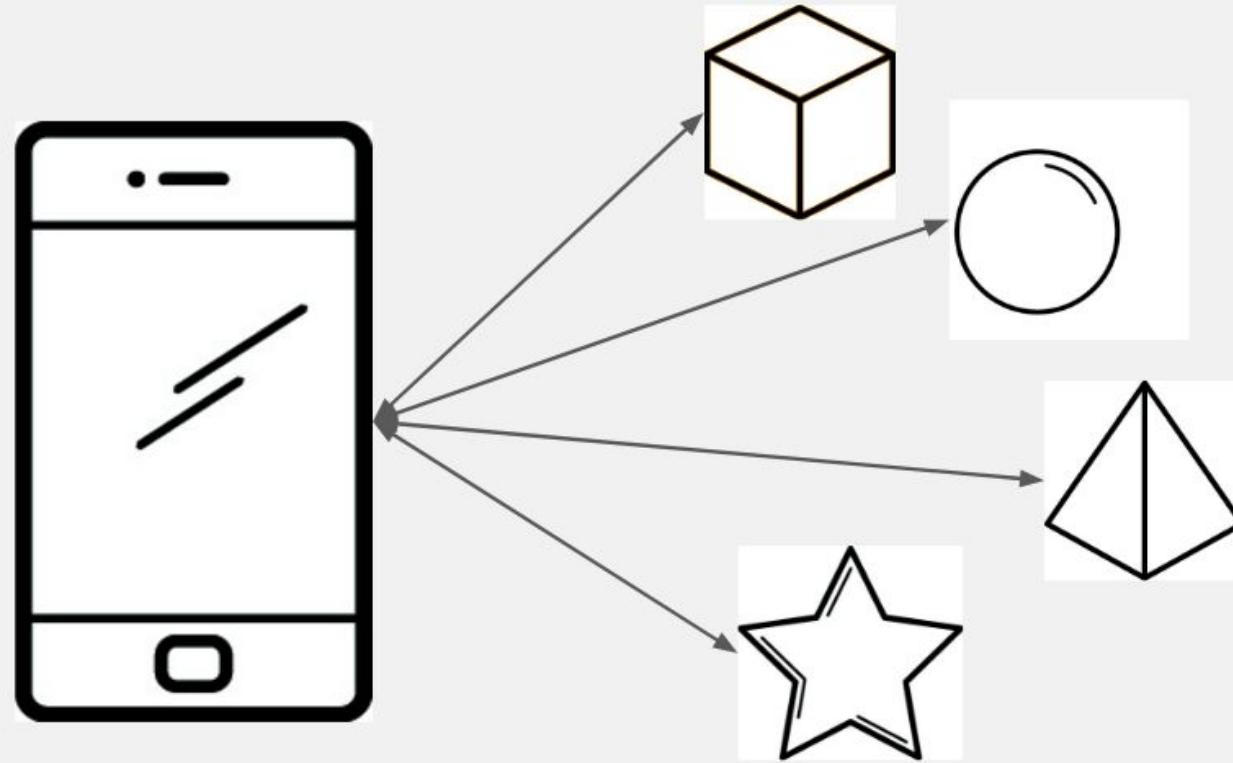


- This is very similar to our CoolStore retail application!
- UI constructed from multiple microservices implemented across business boundaries
- Typically runs in constrained environments like mobile
- Real-world apps consist of 100s of services

# Example



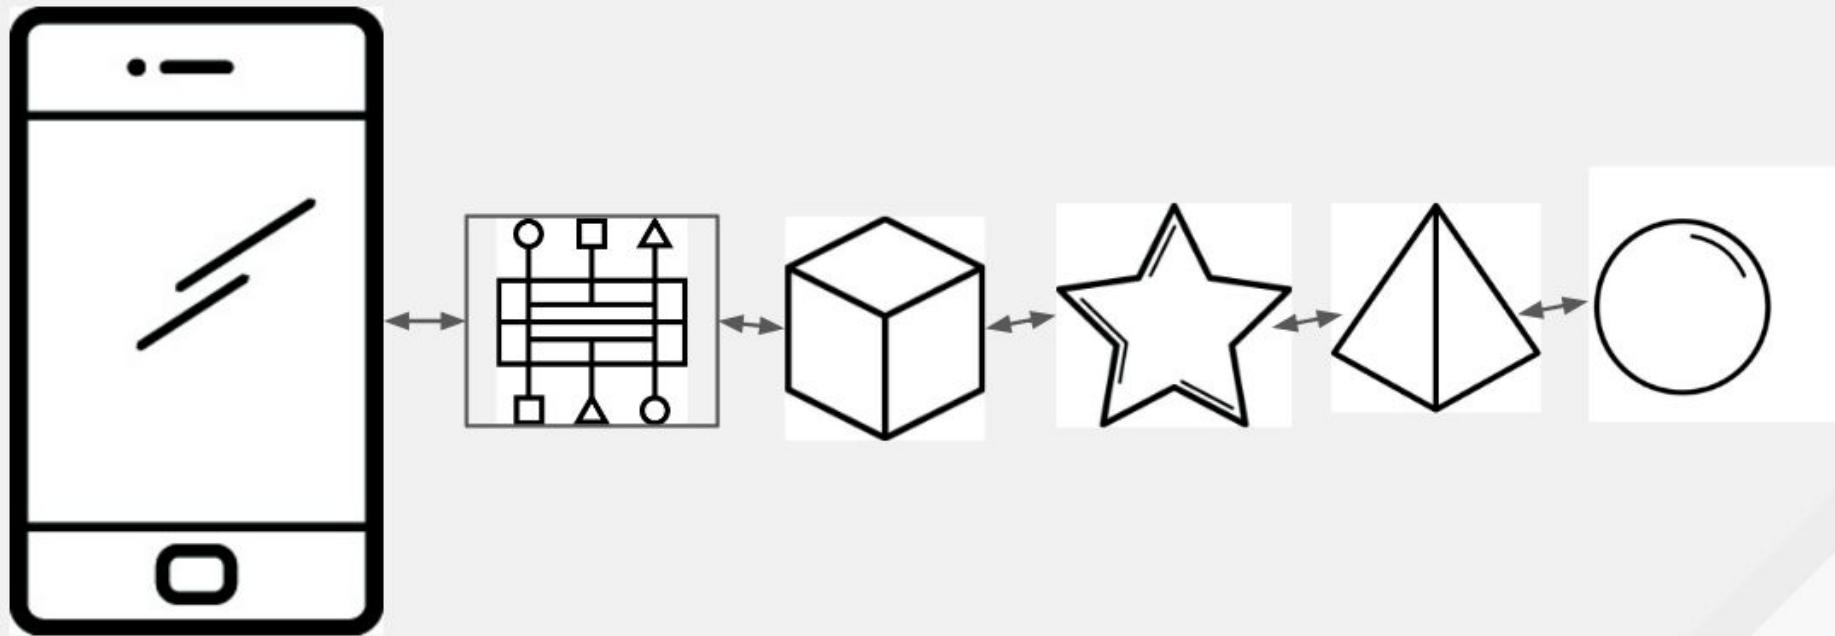
# Option 1: Client Aggregation



# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- Microservices implement functionality previously found in monoliths
- Some microservices depend upon other microservices
- Client applications (e.g. web browsers) depend on *all* of them in one way or another, and are usually “outside the firewall”
- This option means that the client side code (typically run in a browser) is responsible for talking to each microservice and aggregating/combining the results
- Client aggregation benefits
  - No development bottleneck *on the server* / ESB-style funnel
- Client aggregation drawbacks
  - Network bandwidth/latency of multiple calls to multiple microservices
  - Unfriendly protocols - web proxies, ports, etc
  - Difficulty in later refactoring microservices - the client must change too
  - Client application code complexity

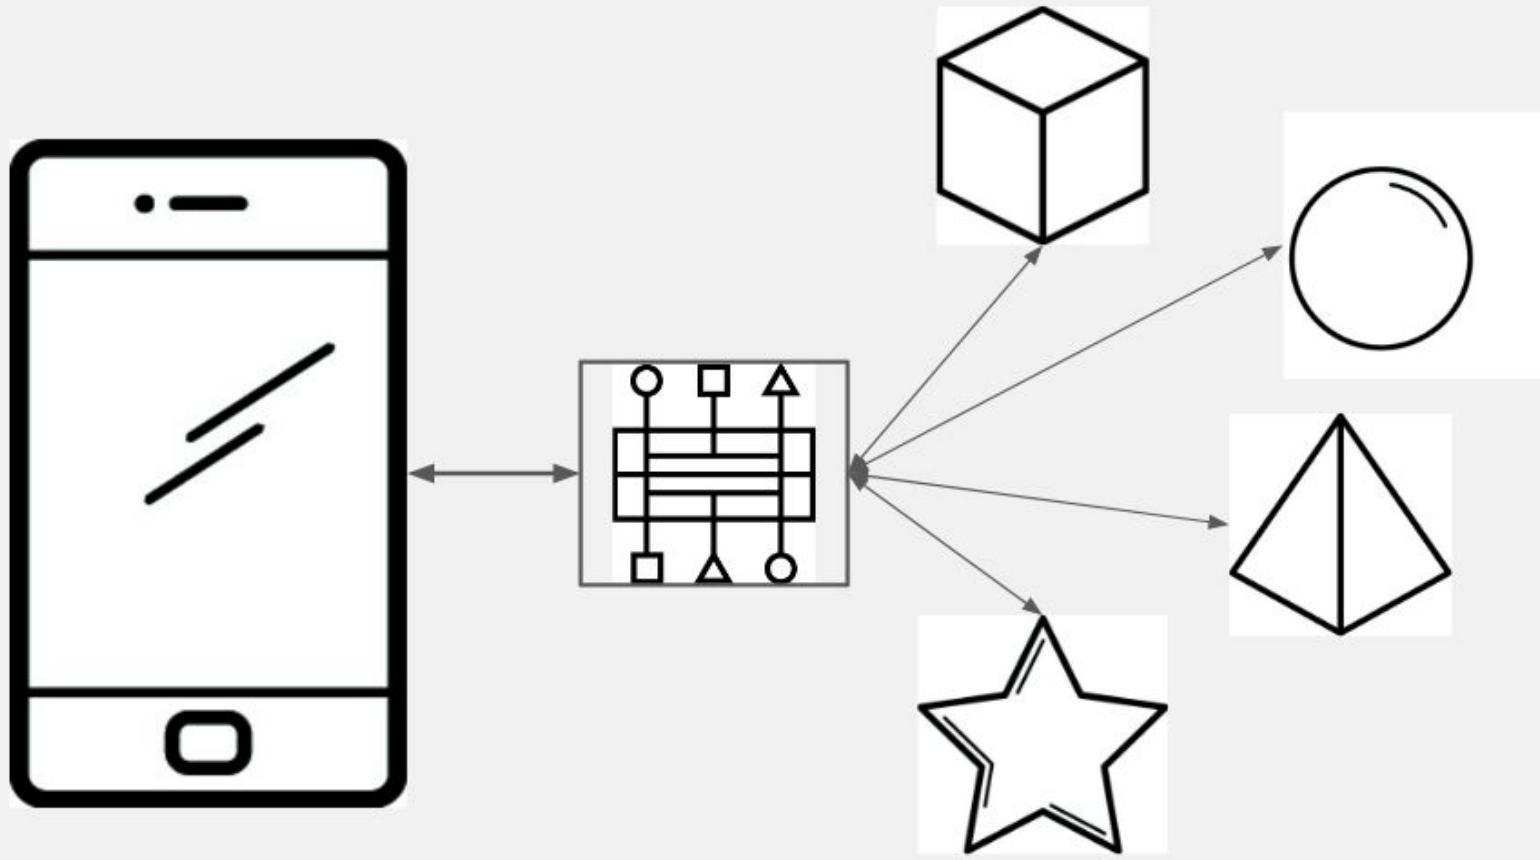
# Option 2: Chaining



# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

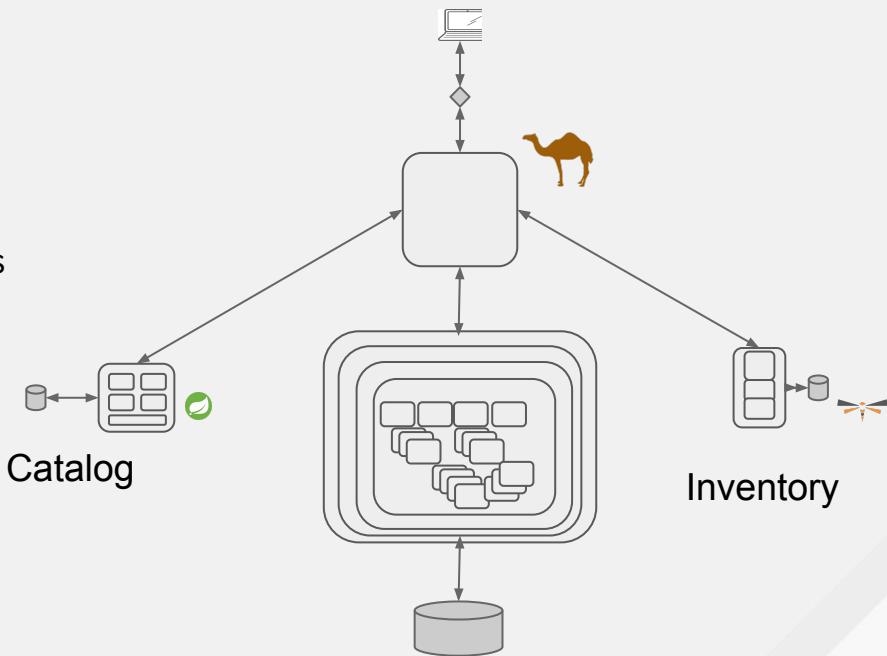
- Chaining means that one microservice calls another, which calls another, etc.
- A complete chain is typically not desirable or necessary, but short chains are OK
- Chaining benefits
  - Client code simpler - there is only a single entry into the chain
  - Less network bandwidth (also due to single entry point)
- Chaining drawbacks
  - Potential for cascading failures (resilience patterns can help minimize this)
  - Complex “stack traces” when things go wrong (tracing libraries a must)
  - Exposes internal structure of app logic (the first microservice in the chain would be difficult to change)

# Option 3: API “Gateway”



# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- The API Gateway pattern:
  - Keeps business logic on server side
  - Aggregates results from back-end services
- API Gateway Pattern benefits
  - Encapsulates internal structure of application's services
  - Less chatty network traffic
  - Simplified client code (no aggregation)
- Drawbacks
  - Possible bottleneck depending on difficulty of adding new services



# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- In this lab, the previously developed microservices will be placed behind a gateway service
- The client application will then call the gateway service to retrieve its data
- This will “strangle” the monolith by replacing its catalog/inventory services with new microservices.
- First, deploy the API gateway:

```
$ cd ~/coolstore
$ git clone -b app-partner https://github.com/modernize-legacy-apps/gateway
$ cd gateway
$ oc project coolstore-<USERNAME>
$ mvn clean fabric8:deploy -Popenshift -DskipTests
```

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- The Coolstore Gateway microservice is a Spring Boot application that implements its logic using an Apache Camel *route*.
- Take a look at the code for the Gateway:
- In your favorite text editor, open  
`src/main/java/com/redhat/coolstore/api_gateway/ProductGateway.java`
- Or in your browser:  
[https://github.com/modernize-legacy-apps/gateway/blob/master/src/main/java/com/redhat/coolstore/api\\_gateway/ProductGateway.java](https://github.com/modernize-legacy-apps/gateway/blob/master/src/main/java/com/redhat/coolstore/api_gateway/ProductGateway.java)

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

```
restConfiguration()
    .contextPath("/services").apiContextPath("/services-docs")
    .apiProperty("host", "")
    .apiProperty("api.title", "CoolStore Gateway API")
    .apiProperty("api.version", "1.0")
    .component("servlet")
    .bindingMode(RestBindingMode.json);
```

- The beginning of this route uses the Camel Java DSL (Domain-Specific Language) to configure the REST system and define the base paths of the API itself (`/services`) and paths to Swagger documentation (`/services-docs`).

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

```
rest("/products").description("Access the CoolStore products and their availability")
    .produces(MediaType.APPLICATION_JSON_VALUE)
```

- This begins the REST DSL portion, defining the primary access point for the catalog of products (/products) and the format of the data it produces (JSON)

```
.get("/").description("Retrieves the product catalog, including inventory availability").outType(Product.class)
    .route().id("productRoute")
        .setBody(simple("null"))
        .removeHeaders("CamelHttp*")
        .recipientList(simple("http4://{{env:CATALOG_ENDPOINT:catalog:8080}}/api/catalog")).end()
        .unmarshal(productFormatter)
        .split(body()).parallelProcessing()
        .enrich("direct:inventory", new InventoryEnricher())
    .end()
.endRest();
```

- This configures the endpoint for retrieving a list of products by first contacting the Catalog microservice, `.split()`ing the resulting list, and *enriching* (via `enrich()`) each of the products with its inventory by passing each product to the `direct:inventory` route.

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

```
from("direct:inventory")
    .id("inventoryRoute")
    .setHeader("itemId", simple("${body.itemId}"))
        .setBody(simple("null"))
        .removeHeaders("CamelHttp*")
    .recipientList(simple("http4://{{env:INVENTORY_ENDPOINT:inventory:8080}}/api/inventory/${header.itemId}")).end()
    .setHeader("CamelJacksonUnmarshalType", simple(Inventory.class.getName()))
    .unmarshal().json(JsonLibrary.Jackson, Inventory.class);
```

- This is the direct:inventory route, which takes in a Product object (in the body()) and calls out to the Inventory microservice to retrieve its inventory. The resulting inventory is placed back into the Camel exchange for enrichment by the enricher:

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

```
@Override  
public Exchange aggregate(Exchange original, Exchange resource) {  
  
    // Add the discovered availability to the product and set it back  
    Product p = original.getIn().getBody(Product.class);  
    Inventory i = resource.getIn().getBody(Inventory.class);  
    p.setQuantity(i.getQuantity());  
    p.setLocation(i.getLocation());  
    p.setLink(i.getLink());  
    original.getOut().setBody(p);  
  
    return original;  
}
```

- This is the enricher logic which takes the Product and matching Inventory objects, *enriches* the Product object with information from the Inventory object, and returns it.
- The resulting list sent back to the client is the list of products, each of which is enriched with inventory information
- The client then renders the aggregate list in the UI.

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- Now that the gateway microservice is deployed, let's hook it into the application using OpenShift routing.
- A route is a way to expose a service by giving it an externally-reachable hostname like [www.example.com](http://www.example.com), or in our example `www-coolstore.<domain>`
- Routes can be created using `oc expose` command line, or through the GUI.
- *Path based* routes specify a path component that can be compared against a URL such that multiple routes can be served using the same underlying service/pod, each with a different path.
- In this case, we already have a route that sends all traffic destined for our monolith to the monolith deployment.
- We want to setup a route such that when the monolith's GUI calls `/services/products`, it is re-routed to our new CoolStore gateway microservice, thus completing the partial strangulation of the Inventory and Product Catalog features of our app.

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- Navigate to Applications → Routes to list the current routes.
- Notice the `www` route is the primary route for our monolith:



- Click *Create Route* to begin creating a new route with the following values:
  - Name: **gateway-redirect**
  - Hostname: The full hostname of the existing route **as seen above** (without the `http://`). For example,
    - `www-coolstore-user21.cloudapps.testworkshop1.openshift.opentlc.com`
  - Path: **/services/products**
  - Service: **gateway**
  - Leave other values as-is (see next page for complete example)
- Click *Create* to create the route

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

## Create Route

Routing is a way to make your application publicly visible.

\* Name

gateway-redirect

A unique name for the route within the project.

Hostname

www-coolstore-user21.cloudapps.testworkshop1.openshift.opentlc.com

Public hostname for the route. If not specified, a hostname is generated. The hostname can't be changed after the route is created.

Path

/services/products

Path that the router watches to route traffic to the service.

\* Service

gateway

Service to route to.

[Split traffic across multiple services](#)

Target Port

80 → 8080 (TCP)

Target port for traffic.

Secure route

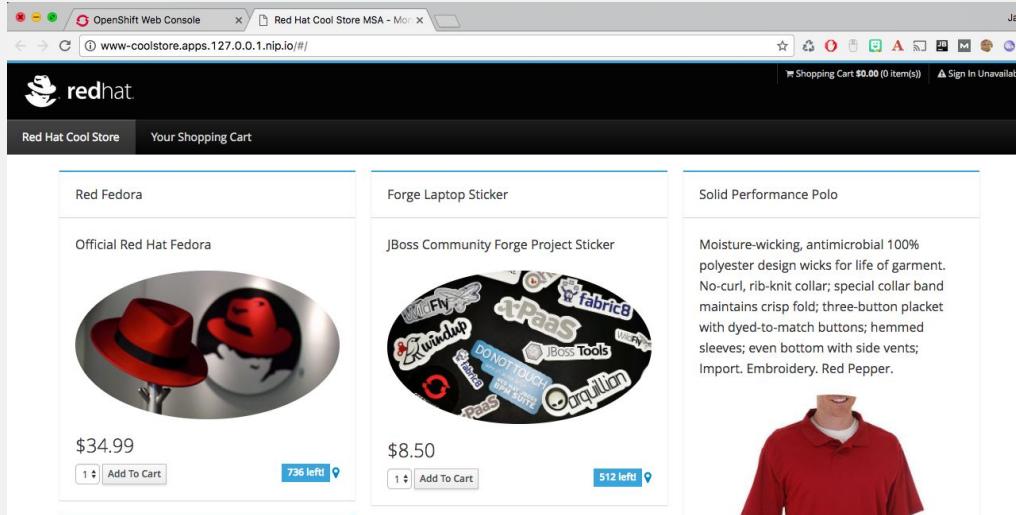
Routes can be secured using several TLS termination types for serving certificates.

**Create**

**Cancel**

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- Test the new route by visiting the application UI (click on the coolstore-prod route in the Overview). It will be no different than the original monolith. How do you know your new microservices are being used in place of the original services?



# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- Let's pretend that there is a lengthy and cumbersome process for getting products and inventories into and out of the backend system.
- We have a high priority task to remove the Red Hat Fedoras from the product list due to a manufacturing defect.
- Let's filter the product out of the result using our gateway. Open the gateway source code file using your preferred text editor (vi, gedit, etc) and un-comment the lines that implement a filter based on product ID (around line 80).
- The highlighted code shows you the predicate used for the filter

`~/coolstore/gateway/src/main/java/com/redhat/coolstore/api_gateway/ProductGateway.java`

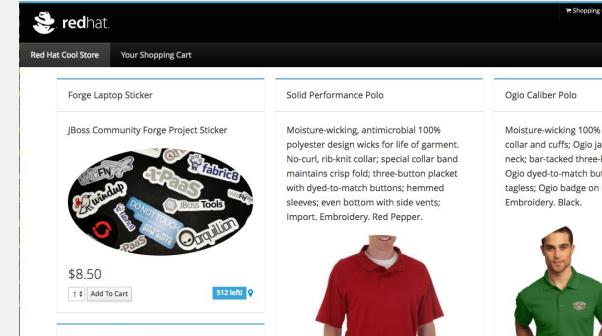
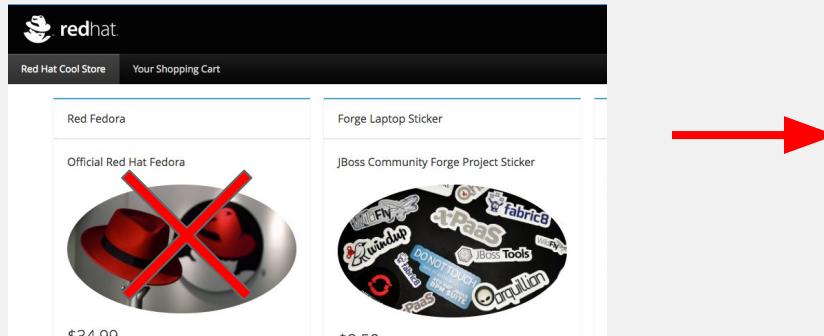
```
//  
// Uncomment the below lines to filter out products  
//  
//     .process(exchange -> {  
//         List<Product> originalProductList = (List<Product>)exchange.getIn().getBody(List.class);  
//         List<Product> newProductList = originalProductList.stream().filter(product ->  
//             !"329299".equals(product.itemId))  
//                 .collect(Collectors.toList());  
//         exchange.getIn().setBody(newProductList);  
//     })
```

# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- Re-deploy the modified code using the same procedure as before:

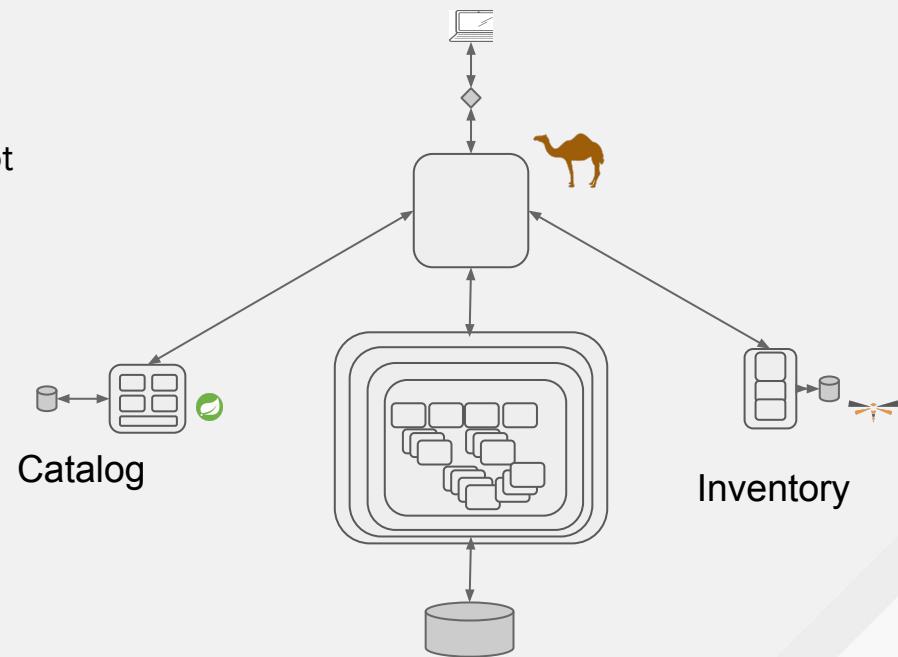
```
$ cd ~/coolstore/gateway  
$ mvn clean fabric8:deploy -Popenshift -DskipTests  
% oc logs -f dc/gateway  
....  
--> Success      # --> wait for it!
```

- Once the new version of the code is deployed and up and running, reload the browser and the Red Hat Fedora product should be gone:



# LAB 3 - MICROSERVICE INTEGRATION PATTERNS

- Our monolith's UI is now talking to our new Camel-based API Gateway to retrieve products and inventories from our WildFly Swarm and Spring Boot microservices!
- Further strangling can eventually eliminate the monolith entirely.





# YOU MADE IT!

James Falkner  
Sr. Technical Marketing Manager, Middleware  
April 2017